Cognitive Mobile Homes

Daniel Greco

Abstract

While recent discussions of contextualism have mostly focused on other issues, some influential early statements of contextualism emphasized the possibility for contextualism to provide an alternative both to coherentism and to traditional versions of foundationalism.[1] In this essay, I'll pick up on this strand of contextualist thought and argue that contextualist versions of foundationalism promise to solve some problems that their non-contextualist cousins cannot. In particular, I'll argue that adopting contextualist versions of foundationalism can both (1) let us reconcile Bayesian accounts of belief updating with a version of the holist claim that all beliefs are defeasible, and (2) let us defend some intuitively plausible epistemological internalist claims from otherwise powerful counterarguments.

Introduction

Epistemic contextualists often motivate their position by arguing that contextualism provides a satisfying resolution of certain skeptical paradoxes.[2] In this essay, however, I will present one version of a very different strategy for motivating contextualism.[3] The strategy I'll explore attempts to motivate contextualism by arguing that it provides an appealing way to chart a middle course between coherentism on the one hand, and traditional, non-contextualist versions of foundationalism on the other.

These two strategies are independent of one another—one can hold that contextualism provides a satisfying resolution of skeptical paradoxes without holding that it

[1] See Williams (1977) and Annis (1978). Arguably, it was also anticipated by Wittgenstein (1969). In his recent book-length defense of contextualism, DeRose (2009) brings up the idea that contextualism has some special relevance to the debate between foundationalism and coherentism, only to set it aside (pp.21-22). I agree with DeRose that there is no entailment from contextualism to either foundationalism or coherentism. But as the rest of this essay should make clear, I do think that that contextualism and foundationalism complement each other nicely.

[2] For paradigm instances of this strategy, see Cohen (1987), DeRose (1995), and Lewis (1996).

[3] Actually, as will become clear later, while I'll focus on contextualism for ease of exposition, my hope is that the strategy can be employed by non-contextualists of various stripes as well. Epistemic relativists (MacFarlane, 2011a,b) and expressivists (Chrisman, 2007) have all offered diagnoses of various phenomena that are structurally similar to contextualist diagnoses. My hope is that any view that can mimic familiar contextualist treatments of, e.g., skeptical paradoxes, (e.g., as the relativist can by appealing to assessment-sensitivity in many of the places where the contextualist would appeal to context-sensitivity) can also mimic the contextualist treatment of foundationalism that I'll offer here, and that the issues that separate contextualism from a variety of rival views will, for the most part, not be relevant to the present discussion.

provides any special advantage to the foundationalist,[4] just as one can endorse contextualist versions of foundationalism without using them to offer the familiar contextualist treatment of skeptical paradoxes.[5] I focus on the less common strategy because I believe it points the way to some surprising further potential applications of contextualism that don't suggest themselves quite as readily when contextualism is motivated via skeptical paradoxes.

1 Foundationalism

I'll use the term "foundationalism" to refer to a view about the structure of epistemic justification. In particular, I'll use it to refer to the view that both of the following claims hold: (1) at least some beliefs are justified even without receiving support from other beliefs—I'll call such beliefs, should any exist, "foundational"—and (2) all justified beliefs are either foundational, or derive all their support (perhaps indirectly) from justified foundational beliefs.[6]

A natural question for the foundationalist is the following: what criteria must beliefs meet in order to be foundational? Historically influential versions of foundationalism held that beliefs must meet very strict criteria in order to be foundational—perhaps they must be impossible to doubt, or must have contents that are of necessity be true whenever believed, or must have some other strong modal properties. It's easy to see why one might find such properties epistemologically interesting. If it were possible to rationally reconstruct our entire body of beliefs from some sparse set of claims that are impossible to doubt, or that must be true (given that we believe them), doing so would be a great intellectual achievement. Criticisms of traditional versions of foundationalism have tended not to target not the desirability of completing such a project, but the feasibility. It's not clear whether *any* beliefs have such strong modal properties, and even if some do (e.g., perhaps the conclusion of Descartes' cogito), they are almost certainly too few to serve as an adequate foundation for the rest of what we normally take ourselves to be justified in believing.

In light of these considerations, contemporary foundationalists tend to defend views on which the requirements for foundational beliefs are more lax, and particular, needn't be free from the threat of doubt.[7] In this first section, however, I want to show how one might be motivated to retreat, not to a laxer invariantist form of foundationalism, but instead to a view according to which the criteria beliefs must meet in order to be

[4] As DeRose does in the passage I refer to in footnote 1.

[5] E.g., while Michael Williams does endorse a sort of contextualist version of foundationalism—he calls it "formally foundationalist" (Forthcoming)—he offers a very different treatment of skepticism from that found in the writings of Cohen, DeRose, and Lewis. In particular, he does not concede that there is *any* context in which global skeptical challenges are successful.

[6] While "foundationalism" is used in too many different ways to speak of a "standard" use, my formulation is similar to what many writers call "minimal" foundationalism (Alston, 1976).

[7] For example, Gilbert Harman (2001) defends a view on which *all* beliefs count as foundational. Phenomenal conservatives defend the only slightly less permissive view that all beliefs that *seem* true to their subjects are foundational for those subjects.

foundational varies with context.

Suppose we're sympathetic to the Cartesian idea that freedom from doubt really is the *sine qua non* for foundational status. After all, we won't be happy treating the claim that P as foundational—justifying other beliefs by showing how they're supported by the claim that P—if we doubt that P. If we allow that which beliefs play the foundational role may vary from context to context, however, then we can insist that foundations must be undoubt*ed*, without taking the Cartesian position that they must be undoubt*able*. We might summarize the two ways of respecting the idea that foundations must be free from doubt as follows:

Cartesian: For a belief to be foundational, it must be undoubt*able*.

Contextualist: For a belief to be foundational (in a context), it must be undoubt*ed* (in that context).

Both the Cartesian and the contextualist agree that our justified beliefs rest on a foundation of undoubted premises. The Cartesian thinks there are certain types of beliefs that can never be in doubt, and so can always play a foundational role. The contextualist disagrees—because she thinks all beliefs are dubitable, she doubts that any beliefs can play a foundational role in *every* context.[8] So while she'll agree with the Cartesian in insisting that foundations be undoubted, the most she hopes for is *contextually local* freedom from doubt, rather than immunity to doubt across all contexts.

A natural thought for the contextualist foundationalist is that when a belief counts as "foundational" in a context, it is treated by the occupants of that context as (at least provisionally) a default assumption in reasoning and argument. Foundationalism is plausible because we must always treat *some* beliefs as default assumptions in reasoning and argument. But it's also plausible that *no* belief will (or should) be treated as a default assumption for reasoning and argument in *every* context. Rather, any belief that we treat as a default assumption in one context, we can treat as a mere hypothesis in another, to be accepted only if it can be supported on independent grounds—i.e., only if it receives support from what we treat as default assumptions in our *new* context.

While the metaphor of Neurath's raft is often appealed to by coherentists, I think it is ultimately most congenial to the contextualist foundationalist. The metaphorical way of putting the point of the previous paragraph is as follows: at any given time, we must stand *somewhere* on the raft, and we cannot examine the planks on which we are currently standing (i.e., each context treats *some* beliefs as foundational, so the coherentist is wrong). But still, we can always shift our weight to new planks, so as to be able to examine the ones on which we were previously standing (i.e., no belief counts as foundational in every context, so the non-contextualist foundationalist is wrong).

There are many ways of developing this idea, not all of which would be labeled "contextualism" in contemporary parlance. The following three views bear a family resemblance to one another, and each of which, in my opinion, can be motivated by the

[8] Or if she grants that *some* beliefs are indubitable, she thinks they are too few to serve the Cartesian's purposes.

considerations I'll discuss in this essay. While I'll generally focus my discussion on what I'll call "orthodox contextualist foundationalism," this is for ease of exposition; most of what I say (with some exceptions that I'll make explicit) should be applicable to each of the views I'm about to mention.

Orthodox Contextualist Foundationalism:

> Sentences of the form "S's belief that P is justified" (as well as sentences attributing foundational justification more narrowly) express different propositions when uttered in different conversational contexts. Features of conversational contexts that affect which propositions such sentences express may include the presuppositions made by the participants to a conversation, the purposes of the conversation, and the practical situation faced by the participants in the conversation. In general, a proposition P's being presupposed by the parties in a conversational context C tends to make sentences of the form "S's belief that P is foundational," true, when uttered in C.[9]

Relativist Foundationalism:

> Sentences of the form "S's belief that P is justified" (as well as sentences attributing foundational justification more narrowly) express propositions that are not absolutely true, but only true relative to contexts of assessment. Which contexts of assessment we take on for the purpose of evaluating sentences about foundationality (and so which such sentences we'll be prepared to endorse) may depend on factors such as which presuppositions we and our interlocutors make, our purposes in evaluating the attribution, and our practical situation (along with, perhaps, the practical situation of the subject of the evaluation).[10]

Expressivist Foundationalism:

> Sentences of the form "S's belief that P is justified" (as well as sentences attributing foundational justification more narrowly) do not express propositions of the usual sort (except perhaps in a deflationist sense), but instead are used to express our acceptance of epistemic norms. As a matter of anthropological fact, which epistemic norms we accept (and so which such sentences we'll be prepared to endorse) varies with factors such as which presuppositions we and our interlocutors make, our purposes in evaluating attributions of foundationality, and our practical situation.[11,12]

[9] E.g., Cohen (1987), DeRose (1995), and Lewis (1996). For the emphasis on presupposition in particular, see Blome-Tillmann (2009).

[10] E.g., MacFarlane (2011b).

[11] E.g., Chrisman (2007).

[12] It may also be the case that epistemic contrastivists such as Schaffer (2004, 2007) can take advantage of some of the arguments I'll discuss.

I have characterized the above views as views concerning the truth conditions of sentences of the form "S's belief that P is justified" or "S's belief that P is foundationally justified." While this might seem like a simple extension of familiar views about sentences of the form "S knows that P", in fact things are not so simple; unlike "knows," "justified" and "foundational" are technical terms in philosophy, rarely used by non-philosophers. While it's easy to see how one might appeal to empirical linguistic data to adjudicate debates about the semantics of "knows", it's harder to see how similar debates can even get off the ground concerning terms that do not occur in ordinary usage. My hope is that while I've expressed the views above as views about the truth conditions of a certain sort of sentence, it's not so hard to see how we can understand them as views about what it takes for beliefs to have certain sorts of epistemic statuses—statuses that *could* be attributed by sentences containing words like "justified" and "foundational," but which can play important theoretical roles even if we don't often directly attribute them. For instance, if I take some belief of yours to have one of the relevant statuses (e.g., to be justified), that will likely have implications for how I interact with you—perhaps I'll be willing to take the belief for granted in the course of planning our joint endeavors, and won't object when you express the belief in assertion, or assert consequences of the belief—even if I never directly attribute that status. And even if we rarely explicitly attribute foundationality (or lack thereof) to beliefs, we may manifest judgments about which beliefs are foundational in other ways. For instance, if I assert "Bob must be in his office," it's plausible that I am not only expressing my belief that Bob is in his office, but also conveying that I do *not* take this belief of mine to be foundationally justified.[13]

There are many important differences (and perhaps some not-so-important ones as well) between orthodox contextualism, relativism, and expressivism, and these differences have been much debated outside the context of their relevance to foundationalism.[14] While the debates have usually concerned knowledge attributions rather than attributions of foundationality; the issues are similar. In fact, given certain additional premises, contextualism, relativism, and expressivism about knowledge *just are* versions of contextualism, relativism, and expressivism about foundations. E.g., Williamson (2000) is naturally interpreted as a foundationalist who holds that a belief that P is foundational for a subject S just in case S knows that P. If we accept this view that knowledge plays this foundational role, then contextualism about knowledge is contextualism about foundations. While Williamson himself has argued against contextualism (2005), the suggestion that Williamsonian views about the centrality of knowledge in epistemology can be fruitfully combined with contextualism is not unheard of.[15]

[13] I am gesturing towards the view that 'must' in English is an 'evidential', in that (one of) the roles of 'must' in sentences of the form "it must be that P" is for the speaker to indicate the nature of her evidence for her belief that P. While the view that 'must' is an evidential is common, it is controversial just what about a speaker's evidence she conveys by using 'must'. See Aikhenvald (2004) for a general discussion of the category of evidentials.

[14] Chrisman (2007) argues that expressivism can reap the benefits of contextualism without some of the main costs. (Though see Plunkett and Sundell (Forthcoming) for a response to this style of argument.) Field (2009) expresses sympathy for both expressivism and relativism in epistemology, and is doubtful that, properly understood, the views ultimately conflict with one another.

[15] See Ichikawa (2013).

My strategy in this paper, then, will be to try to remain as neutral as possible on the issues that separate different versions of "shifty" epistemological views. My hope and expectation is that these differences will not substantially affect how, as a class, these views interact with foundationalism. In particular, I expect that each of these views promises to offer some similar advantages over traditional versions of foundationalism, whatever other drawbacks they may have. For this reason, while I'll focus on arguing that orthodox contextualist foundationalism can address certain problems that non-contextualist versions of foundationalism cannot, I expect that similar arguments could be constructed for the other two views listed above.

So far I've briefly motivated the idea that some version of contextualist foundationalism is a natural response to the infeasibility of the Cartesian project. In the remaining sections of this essay I'll argue that contextualist foundationalism promises to bear some surprising fruit. In the next section, I'll argue that it can resolve a persistent puzzle in Bayesian confirmation theory, and in the section after that I'll argue that it can be used to save an intuitive form of epistemological internalism from otherwise powerful arguments.

2 Bayesianism and Defeat

A number of writers have argued that there is a tension between Bayesian accounts of belief updating, and the holist epistemological claim that all beliefs are defeasible.[16] In this section, I'll first introduce the topic of defeat, and will try to show why, *prima facie*, Bayesian approaches to belief updating promise to deliver an attractive account of how defeat works. I'll then introduce an apparent problem for such approaches—their apparent inability to allow that *all* beliefs are defeasible. While we might respond to this problem by rejecting Bayesian accounts of defeat, or by rejecting the holist claim that all beliefs are defeasible, both of these options involve biting a bullet. I'll argue that we can solve the problem without biting either bullet—at least not exactly—once we view Bayesian theories of belief updating as species of foundationalism, and recognize that this leads to the possibility of a contextualist Bayesian foundationalism. The argument that there *is* a difficult problem for the Bayesian depends, I'll argue, on the assumption that the Bayesian is a non-contextualist foundationalist. If the Bayesian is a contextualist foundationalist, however, the problem has a very natural solution.

2.1 Defeat

Suppose you read in a newspaper that the majority of Salemites intend to vote "Yes" on proposition 17. As a result, you come to believe that proposition 17 will pass. Later, however, you read that the "No" movement is about to release an advertising campaign that, based on testing in focus groups, looks to be extremely effective at changing people's minds.

[16] In particular, David Christensen (1992), Jonathan Weisberg (2009) and James Pryor (Forthcoming).

After you read the first story but before you read the second, you have justification to believe that proposition 17 will pass. But after you read the second story, you no longer have such justification. Epistemologists say that reading the second story "defeats" the support that you had for the claim that proposition 17 will pass.

Can we say anything more precise about defeat? Suppose E is our initial evidence, H is the hypothesis that evidence supports, and E' is our defeater. A natural thought is that in cases of defeat, while E on its own is good evidence for H, E & E' is not.

Moreover, the probability calculus suggests a natural formalization of these relations. After all, probabilistic confirmation is non-monotonic; E can confirm H, even though E & E' does *not* confirm H. Because of the non-monotonicity of confirmation, we can treat cases of defeat as ones in which an initial body of evidence supports some hypothesis, but a larger body of evidence—one that includes a defeater in addition to the initial evidence—does not. Applied to the example just presented, the strategy might go as follows.

> *Intention* = The majority of Salemites intend to vote "Yes" on proposition 17.
>
> *Pass* = Proposition 17 will pass.
>
> *AdCampaign* = A persuasive ad campaign attacking the proposition is about to be released.

Intention is evidence for *Pass*. We can represent this probabilistically by saying that $\mathrm{P}(\textit{Pass} \mid \textit{Intention}) > \mathrm{P}(\textit{Pass})$. The second story is a defeater for the claim that the proposition will pass. If we assume that it the defeater brings us back to roughly the level of support we had initially for *Pass*, we can represent this as follows: $\mathrm{P}(\textit{Pass} \mid \textit{Intention} \ \& \ \textit{AdCampaign}) \approx \mathrm{P}(\textit{Pass})$.

While this is a promising beginning of a Bayesian account of defeat, a full story would have to say much more.[17] The challenge to Bayesian accounts of defeat that I'll discuss, however, challenges any probabilistic account of defeat that turns on the non-monotonicity of probabilistic confirmation, as the proto-account I've sketched so far does.

2.2 The Ubiquity of Defeasibility

The (putative) difficulty for Bayesian accounts of defeat is the following: just as one can undermine support for a hypotheses by attacking the link between one's evidence and the hypothesis, one can also undermine support for hypothesis by attacking the evidence itself. So not all defeat can be understood as involving the accumulation of evidence, where the final, total body of evidence fails to support some hypothesis that was supported by the initial, smaller body of evidence. Arguments of this general form have been given by David Christensen (1992), Jonathan Weisberg (2009), and James

[17] Kotzen (Ms.) does an admirably thorough job of displaying the resources the Bayesian has in describing a wide variety of cases of defeat.

Pryor (Forthcoming). Exactly what this amounts to, and why it's been thought to be a problem for the Bayesian, can be brought out with an elaboration of our earlier example.

As before, you read that the majority of Salemites intend to vote "Yes." But now, the second story doesn't concern an ad campaign, but instead alleges that the polls on which the first article was based come from a source with a flawed methodology—their results are not likely to be representative of what the people of Salem, on the whole, intend.

Intuitively, this defeats your support for the hypothesis that proposition 17 will pass. Why does this present a *prima facie* problem for the Bayesian? Let *Intention* and *Pass* have the same meanings as before, and let *FlawedSource* be understood as follows:

> *FlawedSource* = A second story reported that the polls relied on by the first story come from a source with a flawed methodology.

What we'd *like* to say would be that the following two conditions hold:

1. $\mathrm{P}(\mathit{Pass} \mid \mathit{Intention}) > \mathrm{P}(\mathit{Pass})$
2. $\mathrm{P}(\mathit{Pass} \mid \mathit{Intention} \,\&\, \mathit{FlawedSource}) = \mathrm{P}(\mathit{Pass})$

This would be to treat *FlawedSource* as playing the same role that *AdCampaign* did in our earlier case. The problem is that the second condition *doesn't* hold. If the majority of Salemites intend to vote "Yes", then regardless of how the polls were conducted, the proposition will probably pass. That is, $\mathrm{P}(\mathit{Pass} \mid \mathit{Intention} \,\&\, \mathit{FlawedSource}) > \mathrm{P}(\mathit{Pass})$.

The basic problem is that if we treat the evidence provided by the first story as *Intention*, then that evidence is not defeated by the story about the flawed polling methodology. Rather, if we are to accommodate the possibility of defeat via that story, it looks like we need to treat the evidence provided by the first story as something like *Polls′*:

> *Polls′* = A poll reported that the majority of Salemites intend to vote "Yes.".

Only if something like *Polls* was the evidence that the first story provided can we explain why the story about the flawed methodology defeats your support for thinking that the proposition will pass.[18]

[18] One might think that we could stick with treating your evidence as *Intention* rather than as *Polls* so long as we treat you as Jeffrey conditionalizing on the partition, rather than strictly conditionalizing. Christensen (1992) and Weisberg (2009) convincingly argue that this won't work. While discussing the details is beyond the scope of this paper, the basic problem is as follows. In order to allow that *FlawedSource* undermines the support that *Intention* provides for *Pass*, *FlawedSource* and *Pass* must start out probabilistically independent of one another, but must become probabilistically dependent (in particular, they must become negatively relevant to one another) after Jeffrey conditionalizing on the partition [*Intention*, ∼*Intention*]. But Jeffrey conditionalization can't induce this sort of probabilistic dependence. More intuitively, conditionalizing on *Intention*—whether strictly, or Jeffrey-style—doesn't give us a way to represent that the support for *Intention depends* on beliefs about the polls, and so doesn't give us a way to represent that the support for intention can be defeated by information about polls. For these reasons, I'll ignore Jeffrey conditionalization in the rest of this paper, and will presuppose that we are dealing with strict conditionalization.

So far this might not seem like a problem at all. What's wrong with treating the evidence provided by the first story as *Polls*? The danger is that just as *Intention* can be undermined, so too can *Polls*. For instance, suppose the second story is just a correction—it reports that the first story misreported the poll results. If this defeater is to be accommodated, the evidence will need to concern, not what the polls said, but what the newspaper said the polls said.

Perhaps, at some point, we'll get to a way of characterizing the evidence provided by the first story such that it is immune to the sorts of defeat we've been discussing. Maybe the evidence will concern, e.g., the sorts of experiences I have upon reading it. But to many philosophers, especially those sympathetic to broadly holist epistemological positions, this will seem quaintly optimistic. Even claims about our experiences are not immune to reasonable doubt, and can be defeated in the right circumstances.[19] To hold out hope that there is some class of beliefs that is immune to rational undermining, and is such that all episodes of learning can be understood as involving the acquisition of some new beliefs in the privileged class (with further changes in our body of beliefs coming via conditionalization) is to hope for something very much like the traditional versions of foundationalism that are now generally regarded as failures.

To sum up, the challenge to Bayesian accounts of defeat is as follows. If the Bayesian is too generous about what our evidence is—e.g., if she thinks that our evidence in the cases I've been discussing is *Intention*—then she won't be able to account for certain cases of defeat. But there's no way to avoid being too generous; whatever the Bayesian says our evidence is will turn out to be immune to undermining according to the Bayesian. But nothing is immune to undermining.

In response to considerations like these, some writers have expressed optimism that some alternative approach—perhaps an extension of Bayesianism, perhaps not—will be better able to capture the phenomenon of defeat.[20] To be sure, there are substantial limitations of the Bayesian framework, and if we are to insist on a Bayesian approach to defeat, it should not be on the basis of the view that *all* phenomena of epistemological interest can be captured in the Bayesian framework.[21]

Still, there are reasons to be skeptical of the thought that some alternative approach can provide an improvement over Bayesian accounts of defeat. First, alternative for-

[19]See, e.g., Schwitzgebel (2011) for arguments to the effect that we are highly fallible introspectors. See also Williamson (2000), especially the discussion of our "cognitive homelessness" in chapter 4, for more general arguments to the effect that there is no realm of facts which could play the foundational role that we might have hoped facts about experiences would play. I'll discuss Williamson's position at greater length in the next section of this essay.

[20]Christensen (1992) seems hopeful that some extension of Bayesianism will prove adequate. Pryor (Forthcoming) considers some strategies for handling defeat in a version of the Bayesian framework, but also takes seriously the idea that significant departures from that framework might be required to account for defeat.

[21]To take just one major example, since Bayesianism incorporates an assumption of logical omniscience, it cannot represent logical learning. While there have been attempts to model special cases of logical learning in a Bayesian framework (e.g., Garber 1988), I think it's fair to say that they are not generally regarded as likely to provide an attractive general treatment of logical learning in a Bayesian framework.

malisms available so far run into similar troubles in accounting for undermining defeat.[22] Second, the Bayesian account of defeat handles some cases—e.g., the first version of the newspaper case—extremely nicely, and these cases don't seem all that different from the cases where the Bayesian account seems to falter. If we rest content with the Bayesian treatment of the original case involving a story about polls and a story about an ad campaign, but look for some alternative approach to handle the variations on the original case, we risk ending up with an account of defeat that looks oddly disjunctive. At least to me, it seems that the sort of defeat provided by the story about the ad campaign is not *so* different from the sort of defeat provided by the story about the flawed polling methodology, or the error in reporting the poll results. It would be nice to be able to treat each of these cases as having a similar formal structure, so if we're going to accept a Bayesian account of the first case, (which is hard to resist) it would be nice to have a Bayesian account of the others as well. In the next section, I'll argue that such an account can easily be had, so long as our Bayesianism takes a contextualist form.

2.3 Bayesianism as (Contextualist) Foundationalism

It is instructive to think of Bayesianism as a species of foundationalism, in which a subject's prior probability function, as well as her beliefs in the evidence propositions she updates on, both have a sort of foundational status. Just as non-Bayesian foundationalists hold that a subject's beliefs are justified just in case they derive support from the subject's foundational beliefs, Bayesians hold that a subject's credences are justified just in case they receive the right sort of support from the foundational elements in the Bayesian epistemological framework (i.e., a subject's credences are justified just in case they result from conditionalizing the subject's prior probability function on the subject's evidence propositions).

One reason it is helpful to think of Bayesianism as a version of foundationalism is that it makes clear that the Bayesian faces versions of the question that (as I argued above) makes trouble for non-contextualist foundationalists. Rather than the generic "what criteria must beliefs meet in order to be foundational?" we can pose two Bayesian-specific versions of this question:

1. What criteria must prior probability functions meet in order to be rationally permissible?

2. What criteria must evidence propositions meet in order to be rationally updated upon?

The first question—referred to in the Bayesian literature as the "problem of the priors"—is well known, though it has no generally agreed upon solution. Many historically popular answers to the first question, moreover, can be seen as paralleling historically popular answers to the generic question for foundationalists, and as sharing the difficulties of those answers. For example, Carnap's project of determining a prior

[22] See Weisberg (Forthcoming).

probability function on purely logical grounds shares the appeal, and the infeasibility, of the traditional Cartesian foundationalist project.[23] And the orthodox decision theoretic view that *all* coherent prior probability functions are rationally permissible bears a family resemblance to Harman's view that all beliefs are foundational, and faces similar objections.

The second question—the question about which propositions subjects can rationally update on—is much less often discussed by Bayesians.[24] As the discussion of defeat above should make clear, however, it is not obviously any less difficult than the first; the arguments discussed above seem to suggest that *any* answer to the second question leads to the implausible conclusion that certain beliefs (namely, whichever ones we are entitled to update on) are immune to certain sorts of undermining.

Both of these questions for the Bayesian, however, can be given contextualist answers. And while a contextualist treatment of the first question is a topic for another essay, in what follows I'll argue that a contextualist treatment of the second question can let the Bayesian reconcile her treatment of defeat with a version of the ubiquity of defeasibility.[25]

The basic strategy for saving Bayesian accounts of defeat if we adopt contextualism about what subjects are entitled to conditionalize on, is to argue as follows. First, admit that in a context in which a subject counts as entitled to update on E, E cannot be undermined. So far this looks like biting the bullet and accepting that certain beliefs are immune to undermining—each context will treat certain beliefs as indefeasible. However,

[23]While there are many reasons to doubt that the Carnapian project can be completed, Nelson Goodman's (1983) objections are perhaps the best known.

[24]Deborah Mayo (1996, pp.86-8) identifies the question of when one should accept a piece of evidence—in the Bayesian framework, when one should conditionalize on an evidence claim—as a crucial one for Bayesians, on which they often have little to say.

One notable exception is Timothy Williamson (2000). Like Mayo, he identifies the question as a key one for Bayesians. Rather than rejecting Bayesianism on the grounds that it cannot handle the question, however, he goes on to defend a version of Bayesianism on which a subject's evidential probabilities are determined by conditionalizing a prior probability function on each of the propositions the subject knows.

Still, to the extent that this view does *not* give a broadly Bayesian treatment of knowledge, it can be seen as a sort of retreat, at least if one starts from the point of view of an orthodox Bayesian philosopher of science. As Mayo notes, "the need for a supplementary account of evidence would belie one of the main selling points of the Bayesian approach—that it provides a single, unified account of scientific inference." (p. 88) However, if we're not antecedently committed to Bayesianism answering *all* questions in epistemology and philosophy of science (as I am certainly not) this sort of retreat needn't strike us as unwelcome, even if we find the Bayesian framework a generally useful one.

[25]While I know of no developed contextualist treatment of the problem of the priors, Timothy Williamson seems sympathetic to such a position in his treatment of evidential probability:

> The discussion will assume an initial probability distribution P. P does not represent actual or hypothetical credences. Rather, P measures something like the intrinsic plausibility of hypotheses prior to investigation; this notion of intrinsic plausibility can vary in extension between contexts. (Williamson, 2000, p.211)

Moreover, if Williamson adopted contextualism about knowledge, then his treatment of evidential probability would amount a (broadly) contextualist answer to the second question, since (as mentioned in the previous footnote) he holds that subjects are rationally entitled to update on (and only on) what they know. Though as previously mentioned, Williamson rejects contextualism about knowledge (2005).

the contextualist can account for a version of the ubiquity of defeasibility in the following way: even if we start in a context in which we count as entitled to conditionalize on some evidence E, and in which E thereby counts as immune to undermining, we can always shift to a new context in which we do *not* count as so entitled, and in which our support for E (if any at all) is treated as defeasible, and may be undermined when we learn some new proposition E'.[26] Moreover, it may be that we have a systematic tendency to shift contexts in such a way that claims that in an earlier context are treated as indefeasible are, once the question of their defeasibility is raised, treated as defeasible in a later context. Before applying this strategy to the cases discussed earlier, it may help to consider an analogy to a different sort of contextualism.

Many writers find the view that there are precise cutoffs in the extensions of vague predicates—e.g., that there is a numbers of hairs n such that someone with n hairs is bald, but someone with $n + 1$ hairs is not bald—extremely implausible. But it's very hard to develop plausible alternative theories of how vague predicates work that avoid commitments to precise cutoffs of any sort.[27] One way to blunt some of the counterintuitiveness of precise cutoffs claims is to adopt some form of contextualism about the extensions of vague predicates.[28] Roughly, the idea is that each use of a predicate like "bald" may determine a precise cutoff, even if there is no single cutoff that is determined by all uses of "bald." Vague predicates may have contextually local cutoffs, without having global cutoffs. Stated at this level of generality, it may be hard to see why this should constitute an advance over the more flatfooted theory on which there is a particular cutoff for all uses of "bald." In my view, however, the details of Diana Raffman's (1994; 1996) and Delia Graf Fara's (2000) work show that contextually local cutoffs are significantly more palatable than global cutoffs.

For instance, imagine a particular speaker is faced with a series of increasingly hirsute men, and must classify each as "not bald" or "bald." Just as there is a last straw that breaks the camel's back, there will be a last man in the series that the speaker classifies as bald. But which man this is will be highly sensitive to various aspects of the local context—e.g., depending on factors such as the direction of the forced march (whether the speaker starts with the hairy men or the hairless ones), which men she's recently heard others classify as "bald" or "not bald," and maybe even arbitrary matters like what she had for breakfast, the speaker might draw the cutoff between the bald men and the men who are not bald differently. If we think that this sort of usage is correct—that such a speaker would be using "bald" just as she should—then it's tempting to adopt a theory on which uses of vague predicates *do* determine precise cutoffs, but that

[26] Ram Neta (2003, 2004, 2005), has defended contextualism about evidence, and the view I express here is heavily indebted to his work. To my knowledge, however, he has never combined his contextualism with Bayesian approaches to belief updating, or used it to address the apparent challenge to such puzzles presented by the ubiquity of defeasibility. The strategy I pursue in this section also bears some similarity to Lewis' (1996) contextualist strategy for saving closure principles for knowledge from the threat of Kripke's dogmatism paradox. Like Neta, though Lewis does not apply his discussion of contextualism to Bayesianism.

[27] See Williamson (1994).

[28] See, e.g., Raffman (1994, 1996), Fara (2000).

these cutoffs are sensitive to the idiosyncratic dispositions and interests of particular speakers on particular occasions of utterance.[29] If this theory can be supplemented with an explanation of why the cutoffs never seem to be right where we look—e.g., why it never seems that for some particular pair of men with n and $n+1$ hairs, the first is bald and the second is not—then the counterintuitiveness of the claims it makes about precise cutoffs will be further blunted. And both Raffman and Fara provide such explanations.

The analogy to the present discussion works as follows. The contextualist views about vagueness I've just been discussing have two main components:

1. An account of how particular utterances of vague predicates determine precise cutoffs, and

2. An explanation of why the cutoffs never seem to be where we're looking (and so, derivatively, an explanation of our temptation to deny claims about precise cutoffs)

The strategy that I'm suggesting for the Bayesian will need to have two analogous components:

1. An account of how particular contexts do treat certain beliefs as indefeasible, and

2. An explanation of why no particular belief we consider ever seems to be indefeasible (and so, derivatively, an explanation of our temptation to endorse the claim that all beliefs are defeasible)

While developing these components in the level of detail that Raffman, Fara, and others have done in the case of contextualism about vagueness is well beyond the scope of this paper, I hope the following will give the reader some idea of how the story might go, as applied to the cases of defeat discussed earlier in this section.

As for the first component of our strategy, the short answer is that in any context, certain claims will count as legitimate to update on, and for reasons we've already seen, what you update on will be immune to undermining, in the Bayesian framework. For instance, we can agree in the initial case that you are entitled to conditionalize on *Intention*, and that doing so provides support for *Pass*, but that this support is defeated when you learn *AdCampaign* The initial, flatfooted Bayesian account of what's going on when you read one story about what the voters intend, and then another story about how those intentions are likely to change, can remain untouched.

What about the next case, where the second story concerns flawed polling methodology? If we started out discussing the earlier case, then bringing up this more complicated one is liable to induce a shift into a new context in which you no longer count as having been entitled to conditionalize on *Intention* in the first place. In this new context—one in which we are taking seriously possibilities in which the polls are unreliable—all you're entitled to conditionalize on is *Polls*—the claim that the polls reported that the majority of Salemites intend to vote "Yes." While this claim on its own supports *Pass*, together

[29] Raffman (1994) focuses on the forced march version of the sorites, and on the idea that correct usage will involve making arbitrary distinctions that differ from one occasion of use to another.

with *FlawedSource*—which we still take you to be entitled to conditionalize on—it does not.

The next step is similar. When we move to the case in which the second story reports that the first story misreported the poll results, we induce yet *another* context shift. Now that we don't regard you as having unproblematic access to what the polls said, we treat you as not entitled to conditionalize on claims about what the polls reported. Rather, we'll think you're only entitled to conditionalize on claims about what the newspaper *said* the polls said, and this claim won't support the hypothesis that the proposition will pass when combined with the evidence you get from the second story.

Whenever we regard a claim as having been defeated, we'll give the same formal account; we'll regard somebody as having had an initial body of evidence E that supported some hypothesis H, but as having a final body of evidence E'—a superset of E—that no longer supports E. That is, defeat will always be explained as involving only evidence *accretion*, and never evidence *deletion*. The role that context plays, however, is in determining what gets to count as the initial evidence; rather than giving a context-independent, univocal answer to the question of what a subject's evidence is, we allow that our answer to questions about what a subject's evidence is will change, as our context changes (e.g., we may initially regard you as entitled to update on claims about what the majority of Salemites intend, but we also recognize that if our context shifts, we may no longer regard you as so entitled).

Just what induces these context shifts? This is closely related to the demand for the second component of our strategy—the demand for an explanation of why no particular belief we consider ever seems to be indefeasible. The answer will depend on the particular form our contextualism takes, and things will of course look different if we instead opt for some version of relativism or expressivism. But certain basic factors may be common to each approach in the family—perhaps subjects will generally not count as entitled to update on some proposition when there are salient possibilities (salient to whom? that may differ depending on the approach) in which they are in error as to the truth of that proposition. Most of the time we're ignoring lots of error possibilities, and we'll treat subjects as entitled to conditionalize on claims about their external environment. But in special cases where certain error possibilities are raised to salience, we'll retreat, and treat subjects as entitled to conditionalize only on other claims—perhaps claims about their experiences, or claims about their inclinations to believe claims about their experiences, or something else. And there isn't any principled stopping point to this process of retreat—there's no sort of evidence statement about which no doubts could be raised, though in practice we're usually not inclined to retreat very far.[30]

We can use these observations to provide the basis for a general story about why no beliefs seem to be indefeasible; raising the question as to whether some belief is defeasible essentially amounts to an invitation to consider possibilities in which the belief might be false, so once the indefeasibility of a belief is in question, the context will typically shift so as to ensure that the belief is defeasible. That is, once we are considering possibilities

[30]Moreover, the "retreat" needn't always be to weaker and weaker propositions. Rather, as the Neurath's raft metaphor suggests, it might simply be to *different* propositions.

in which we are in error as to whether E, it will seem irresponsible to simply take E for granted; rather, we'll be inclined to accept E only if it receives adequate support from other claims we *do* take for granted. But this is just what it is for us to treat E as enjoying merely defeasible support. Just as—at least according to Raffman and Fara—raising the question of whether the cutoff for "bald" lies between a particular pair of adjacent men in a sorites series will typically bring us into a new context in which the cutoff does *not* lie between those two men, even if it did in our old context, raising the question of whether some claim E enjoys indefeasible support will typically bring us into a new context in which E enjoys merely defeasible support, even if it initially enjoyed indefeasible support in our old context.

The apparent difficulty for Bayesian accounts of defeat required assuming that the Bayesian must give a single, context-independent answer to the question of what you should conditionalize on when you read a particular newspaper story. Whatever answer she gives, as we saw before, she's in trouble. But now we see that the Bayesian can reject the demand to tell us, once and for all, what you should conditionalize on in such a case. In most contexts, where we're not worried about certain error possibilities, there's no obstacle to allowing that you're entitled to conditionalize on the content of what the article reports (rather than just, e.g., the claim *that* it reported what it did). But other contexts will demand other treatments, and we needn't allow that there's some privileged, stable context that we could never be pushed out of—we needn't allow that there's any answer to the question of what you should conditionalize on that we couldn't come to reject after yet another context shift.[31] The resulting position can be summed up, more or less, by the following remark from Wittgenstein's *On Certainty*: "the same proposition may get treated at one time as something to test by experience, at another as a rule of testing." (1969, 98)

In the absence of more general motivations for contextualist foundationalism, this reply on behalf of the Bayesian might seem ad hoc. But once we admit both that (a) contextualist foundationalism is a viable contender in debates about the structure of epistemic justification, and that (b) Bayesianism is in effect a species of foundationalism

[31]Earlier in this section I suggested that it would be a defect of a treatment of defeat if it were "oddly disjunctive." But one might worry that the sort of account I've suggested on behalf of the Bayesian is just that. We might put the worry as follows: on the story I've been sketching, there are really two quite different types of defeat. The first is the simple sort of defeat that doesn't involve anything like context shifting.

This charge, however, can be resisted. The sense in which the contextualist can deny that she countenances two forms of update (or two associated forms of defeat) parallels the contextualist strategy for affirming "intellectualism"—the thesis that what one knows depends only on "truth-conducive factors." Roughly, while the contextualist will accept different sentences involving epistemological vocabulary (whether those are knowledge attributions, or claims about what subjects are entitled to update on) as factors other than truth-conducive factors change, (perhaps as the practical stakes change, or as she starts considering new error possibilities) in no context will she affirm that some subject would (or would not) know/justifiably believe some proposition if only the non-truth-conducive factors were different (e.g., if only the practical stakes were different, or if only we were considering fewer error possibilities). See Stanley (2005, pp. 2-3) and DeRose (2009, pp. 24-5), both of whom accept that the sensitive invariantist must deny intellectualism, while the contextualist may accept it. Moreover, I take it that this feature of contextualism is shared by expressivism and relativism, for essentially the same reasons.

in which the beliefs that subjects are entitled to update on play a foundational role, it's hard to avoid the conclusion that (c) we must take seriously contextualist Bayesian foundationalism, according to which the beliefs that subjects are entitled to update on play a foundational role, but which sorts of beliefs get to play that role varies with context. And once we take that view seriously, we see that it has the considerable virtue of providing the foundations for a general account of defeat.

Still, it's worth acknowledging that the sort of contextualism we must appeal to in order to employ the strategy I've been discussing is, in important ways, a departure from some familiar forms of contextualism. Stalnaker (2008, pp.102-105) distinguishes between *superficial* and *deep* contextualism. Roughly, the sort of context-dependence posited by the superficial contextualist is eliminable, while the sort of context-dependence posited by the deep contextualist is not. The deep contextualist is contextualist "all the way down," while the superficial contextualist thinks that, once you start digging, you eventually hit some context-independent epistemological bedrock. To illustrate the distinction, it will help to have some examples.

Consider the following simple form of contextualism about justified belief. A subject's belief that P is properly called "justified" in a given context only if the epistemic probability of P for the subject is above a certain threshold. Which threshold? That depends on the context. On this view, the epistemic probability of P does *not* depend on context—context only comes into the picture in determining how probable P must be in order for a subject's belief that P to merit the label "justified."[32] This is a version of superficial contextualism, since it allows an epistemological notion—epistemic probability—for which contextualism does not hold.

David Lewis (1996) defends another version of superficial contextualism. According to Lewis, knowing that P requires eliminating all contextually relevant counterpossibilities to P. His view is contextualist because the set of relevant counterpossibilities varies with context. But what does the eliminating—what one's evidence is—does not.[33] On Lewis' view, one's evidence is always given by the content of one's experience and one's memory (1996, p. 425). If we put things in Bayesian terms, this amounts to the claim that one is always entitled to update on the content of one's experience and memory, no more and no less. For my purposes, this is a crucial juncture at which Lewis' contextualism takes a superficial rather than a deep form, crucial because the Lewisian contextualist cannot appeal to my strategy for reconciling Bayesian accounts of updating with the apparent ubiquity of defeasibility; Lewis' view straightforwardly implies that our knowledge of the contents of our experience and memory is indefeasible in every context, so if we want to reject that assumption we'll need to look elsewhere.

[32]The view defended by Stewart Cohen (1988) isn't *too* far from this one. He advances a contextualist version of the relevant alternatives theory of knowledge; on his view epistemic probability is taken as fixed independently of context, and how probable an alternative has to be before it counts as relevant is determined by context. Also, while they're sensitive invariantists rather than contextualists, Fantl and McGrath (2009) employ a similar strategy on which epistemic probability is stakes-invariant, but the cutoffs for knowledge and justification are not.

[33]Stalnaker (2008, pp.102-5) gives further reasons why Lewis' view is best thought of a superficial form of contextualism.

Most writers who've defended versions of contextualism have only officially committed themselves to superficial contextualism, since they've only defended contextualism about one or another epistemological notion. Moreover, insofar as they invoke some context *in*sensitive epistemological machinery in explaining their position—e.g., a context insensitive notion of epistemic probability—they're committed to rejecting deep contextualism.[34] Still, while deep contextualism is a radical thesis, I think it deserves to be taken seriously.[35] I've already shown how deep contextualism (or at least, a form of contextualism that's deeper than some extant popular forms of contextualism) can help the Bayesian. In the next section of this essay, I'll point to another unobvious virtue of contextualist foundationalism. While it's at least not *obvious* to me that only deep versions of contextualist foundationalism have this virtue, my suspicion is that, as in the case of Bayesianism and defeat, superficial contextualists will not be able to avail themselves of the argumentative strategy I suggest.

3 Internalism and Mobile Homes

In this section I'll argue that contextualist foundationalism provides us the necessary resources to rescue a *prima facie* attractive version of access internalism from otherwise powerful objections.

3.1 Motivating Access

Consider the following strong "access internalist"[36] epistemological principle:

(Access) If S ought to believe that P, then S is in a position to know that S ought to believe that P.[37]

There are various ways of providing some *prima facie* motivation for this principle. I'll consider just two. The first turns on an analogy to ethics. Many ethicists distinguish

[34] As an aside, I'm inclined to think that many versions of superficial contextualism aren't particularly interesting qua epistemological theses (though they may be interesting qua semantic theses); to the extent that the contextualist allows a realm of "pure," context-insensitive epistemology, epistemologists can largely ignore context and focus their investigations on topics to which it isn't relevant. Moreover, the existence of a set of context-insensitive epistemological questions limits the scope of the epistemological work that can be done by contextualist theses. For instance, if we try to solve skeptical problems by appeal to some version of superficial contextualism, such problems will likely reappear concerning epistemological notions where we concede that contextualism won't help. If we are contextualists about knowledge but not epistemic probability, then even if our contextualism helps us explain why utterances of "I know that I have hands" might be true in ordinary contexts, it won't help us explain why the epistemic probability that I have hands is any higher than the epistemic probability that I'm a handless brain in a vat.

[35] See Stalnaker (2008) for further motivation.

[36] For a discussion of access internalism, and how it differs from other versions of internalism, see Conee and Feldman (2001).

[37] While I focused on outright beliefs in this paper, there are probabilistic versions of (Access) as well, and the strategy I provide here applies, *mutatis mutandis*, to the claim that subjects are always in a position to know which credence functions they have justification to adopt.

between an "objective" sense of ought claims, and a "subjective" sense of ought claims, where roughly, what I objectively ought to do is what there is most reason to do in light of *all* the information, and what I subjectively ought to do is what there is most reason to do in light of the information available to me.[38] For instance, suppose my friend is hungry, and I have in my possession some bread that is, unbeknownst to me, poisoned, but which I have every reason to believe is healthy and nourishing. Ethicists will say that I objectively ought not give my friend the bread, but I subjectively ought to give her the bread.

One reason many ethicists are interested in the subjective ought is that the subjective ought is thought to be "action-guiding" in a way that the objective ought is not.[39] It's very hard to make out just what this sense of "action-guiding" amounts to, but it's tempting to try to explain it by appeal to the idea that subjects are often not in a position to know what they objectively ought to do, but they are always in a position to know what they subjectively ought to do. That certainly seems to be what's going on in the example of the poisoned bread—I can't know that things wouldn't go best were I to give my friend the bread, but I *can* know that giving my friend the bread is the best option in light of the information available to me.

But the claim that we are always in a position to know what we subjectively ought to do is an ethical analogue of (Access). That is, we can draw a similar distinction between objective and subjective oughts in epistemology, and the same sorts of motivations that push ethicists to think that, if the subjective ought is to be action-guiding, we must always be in a position to know what we subjectively ought to do, will push us to accept (Access). It's natural to think that insofar as there is an epistemic analogue of the "objective ought," we always epistemically ought to believe the truth, in this objective sense of the epistemic ought.[40] But that's not the sense of epistemic justification, or epistemic oughts, that epistemologists are interested in. We might try to explain why that is by saying that we are not always in a position to know what is true, so the objective epistemic ought is not "deliberation-guiding," but we *are* always in a position to know what we have justification to believe, so epistemic justification *is* deliberation-guiding.

Claims like this are apt to raise suspicions, in part because there are powerful reasons to doubt that there could be *any* sort of claim that we are always in a position to know whenever it is true, so we might worry that the relevant sense of "action-guiding-ness" that the subjective ought is supposed to have, and "deliberation-guiding-ness" that we might hope epistemic justification could have, is too strong for any interesting normative notion to satisfy. (Access) just amounts to the claim that the condition of having justification to believe that P is "luminous" in the sense of luminosity for which Williamson

[38] A common view in recent years involves holding that "ought" isn't really two-way ambiguous or polysemous, but rather that "ought" claims must be understood as involving a kind of tacit relativity to a body of information, which might be all the information (as in the "objective ought"), might be the subject's information (as in the "subjective ought"), but might be some other body of information too. See Finlay and Björnsson (2010), and Kolodny and MacFarlane (2010).

[39] See, e.g., Hedden (2012), Sepielli (Forthcoming).

[40] See Gibbard (2005).

(2000) argues that *no* non-trivial conditions are luminous. I'll discuss worries in this neighborhood soon—for now I'm just trying to explain why one might be attracted to (Access). Later in this section I'll try to show that adopting contextualist foundationalism can let defenders of (Access) avoid objections in the spirit of Williamson's.

The second strategy for motivating (Access) that I'll discuss involves focusing on the impropriety of various broadly Moore-paradoxical claims, and arguing that this impropriety is best explained by (Access). Consider claims like the following:

1. P and I do not have justification to believe that P.
2. I have justification to believe that P and it is not the case that P
3. P and it is an open question whether or not I have justification to believe that P
4. I have justification to believe that P and it is an open question whether or not P.[41]

If (Access) were true, it would provide an easy explanation of the impropriety of 1-4. With slight modifications depending on which claim's impropriety we're explaining, the general strategy is as follows. Whenever a subject has justification to believe P, if (Access) is true, then the subject is in a position to know *that* she has justification to believe that P. So if she believes P, but fails to believe that she has justification to believe P (either because she believes that she lacks justification to believe that P, or because she regards it as an open question whether she has such justification), then she is going wrong in one of two ways. Either (1) she lacks justification to believe P (and so is unreasonable in believing it), or (2) she has justification to believe P, but she is failing to believe something that she is in a position to know (namely, *that* she has justification to believe that P). If failing to believe something that one is in a position to know is a mistake, from an epistemic point of view, then we can explain why somebody who believes one of 1-4 is always makings some kind of epistemic error.

Providing such explanations is a good deal harder if we don't accept anything like (Access)—3 and 4 are particularly tricky. If (Access) is false, then there are cases in which one has justification to believe some proposition P, but one is not in a position to know that one has such justification. But then it seems like the appropriate response to such cases should be to believe a claim like 3—to believe that P, while regarding it as an open question whether one has justification to believe that P. So if believing claims of the form of 3 is always improper, then we have some motivation to accept (Access).[42]

[41] This list, and the strategy of motivating claims like (Access) by appeal to the impropriety of claims like the ones on the list, is from Smithies (2012). See also Feldman (2005), who pursues a more limited version of a similar strategy, and Horowitz (Forthcoming), who focuses on the impropriety of claims like 1, and in particular how believing them can lead to unreasonable behavior as well as unreasonable inferences as more and more such claims are believed.

[42] If it doesn't strike the reader as implausible that believing such claims is always improper, see Smithies (2012) and Greco (2014b) for some persuasion. See also Sosa (2009), who argues for the impropriety of some similar claims, though does not take such impropriety to support (Access).

3.2 Access and Traditional Foundationalism

Even if strong *prima facie* motivations for (Access) can be assembled, there are powerful reasons to think that nothing of the sort could be true. Suppose (Access) *were* true. What could explain its truth? A natural, almost inescapable thought, is the following. (Access) could only be true if there were some distinctive realm of propositions that constituted the supervenience base for facts about what we have justification to believe, such that concerning claims in this distinctive realm, one is always in a position to know the truth.[43] While this might not be *enough* to guarantee the truth of (Access), it seems at least required. After all, if the supervenience base for facts about justification is such that in some cases we are *not* in a position to know the truth concerning it, then such cases will be natural candidates for cases in which we are also not in a position to know what we have justification to believe.

In light of such considerations, some defenders of (Access) do endorse strong claims about our epistemic access to the facts on which justification supervenes; Smithies (2012) defends a version of (Access), and holds that the truth of his version of (Access) is explained by the fact that facts about what we are justified in believing supervene on facts about a very special class of mental states. According to Smithies, "The determinants of justification are non-epistemic facts about one's mental states, which are introspectively accessible in the sense that one has introspective justification to believe that those mental facts obtain if and only if they obtain."[44] (2012, p.297) Smithies goes on to hold that since it is an *a priori* matter which beliefs are justified by which introspectively accessible mental states, one is always in a position to know what one has justification to believe, by a combination of introspection (to find out which introspectively accessible mental states one is in), and *a priori* reflection (to find out what beliefs those mental states justify).

If this is the only sort of strategy available for explaining (Access), however, many philosophers will be inclined to deny the truth of (Access), rather than to accept an inference to the best explanation argument for the existence of a distinctive class of "non-epistemic facts about one's mental states," with the feature that one "has introspective justification to believe that those mental facts obtain if and only if they obtain." After all, to accept the existence of such a class of facts is to accept something very much like a traditional, Cartesian sort of foundationalism. It is to accept that we have what Williamson (2000) has called a "cognitive home," a realm in which nothing is hidden to us—a class of claims such that, if we pay close enough attention (for Smithies, as well as for Descartes, this will involve doing careful enough introspection), we will always know which such claims are true, and which are false. Moreover, it is to hold that this cognitive

[43] E.g., Srinivasan (Ms.) presupposes that (Access) requires a Cartesian view of the mind, where the supervenience base for facts about epistemic obligation is some set of facts about our mental states, which we are always in a position to know. And this view *is* held by the most explicit contemporary defenders of (Access), as I discuss in the next paragraph.

[44] One difference between my formulation and Smithies' is that while I've put (Access) in terms of being in a position to know what one has justification to believe, the claims he discusses are put in terms of having justification to believe the truth about what one has justification to believe. I don't think this will make a difference for my purposes.

home isn't trivially or uninterestingly small—it doesn't only include, e.g., claims to the effect that we exist—but instead it includes enough to form a supervenience base for all facts about what we have justification to believe.

But there are both general theoretical reasons to doubt that we have any such cognitive home, as well as specific reasons to doubt that claims about our mental states in particular could constitute such a home.[45] Rather than try to rebut these claims, in the next section I'll try to argue that explaining (Access) does not require positing the existence of a cognitive home, or at least not one of a traditional sort. Rather, if we adopt contextualist foundationalism, we can explain (Access) by appeal to a sort of cognitive mobile home—in each context, we'll regard some class of facts as constituting a subject's cognitive home, but when reasonable doubts about the subject's ability to know facts in the relevant class are raised and taken seriously, the context will shift, and so will the class of facts that we take to constitute the subject's cognitive home.

3.3 Mobile Homes

The basic strategy that the contextualist can use to defend (Access) is to hold that contexts determine *local* supervenience bases for facts about justification, rather than global ones.[46] Moreover, she can hold that it is a constraint on B's counting as a local supervenience base for facts about what a subject S is justified in believing, that in the context in which B counts as such a supervenience base, S counts as in a position to know the truth concerning B. To see what this amounts to, and how it can help the defender of (Access), it will help to consider an example.

Suppose we're in a context in which the following facts are all taken for granted: Peter is hosting a party, and he has either said to Justine that the party starts at 7:00, or that it starts at 8:00. Whatever Peter said, Justine heard him.

Now someone raises the question of what Justine ought to believe concerning the time at which the party starts. A natural answer is that Justine should believe whatever Peter told her—if Peter said the party is at 7:00, that's what Justine has justification to believe. *Mutatis mutandis* if he said the party is at 8:00. That is, the facts about what Justine ought to believe locally supervene (given the background presuppositions of our context) on the facts about what Peter said. And because we take it that Justine heard Peter—she knows what Peter said—we'll take it that Justine is in a position to know what she ought to believe.

[45] Quine (1951/1953) is probably the most influential critic of the general picture on which all of our knowledge rests on beliefs about matters in a "cognitive home." For some reasons to doubt that facts about our phenomenology in particular constitute such a home, see Schwitzgebel (2011).

[46] The distinction between local and global supervenience bases should become clear as this section progresses. As an aside, it needn't be an idiosyncratic feature of facts about justification that they only have local supervenience bases. See Rayo (Forthcoming) for a "localist" view of representation on which, quite generally, contexts only determine local supervenience bases for facts of most sorts. It is the sort of distinction between "local" and "global" supervenience bases that he is talking about that I mean to be tracking with my terminology; the fact that philosophers of mind use the words "local" and "global" to distinguish between different versions of supervenience is, from the standpoint of the present paper, an unfortunate coincidence—I do not mean to be tracking the same distinction that they are.

To be sure, facts about what Peter said would not make a plausible *global* supervenience base for facts about what Justine has justification to believe concerning the start time of the party. There are certainly contexts in which we would *not* hold that what Peter said fixes what Justine ought to believe concerning when the party starts. One obvious sort of situation in which we might cease regarding facts about what Peter said as such a supervenience base would be if we were to come to doubt that such facts provide good evidence concerning the start time of the party—if, e.g., Peter had a reputation as a Prankster, we might not think that Justine ought to trust him. But another reason we might cease regarding facts about what Peter said as a supervenience base for Justine's justification would be if we were to come to doubt that Justine was in a position to know what Peter said.

Suppose, e.g., we come to think that Peter might've mumbled, or that Justine has trouble hearing. If we continued to hold that Justine ought to believe whatever Peter said, then we would have to regard Justine as not being in a position to know what she ought to believe. But a more likely response, I think, is that we will change our minds about which facts determine what Justine ought to believe. Perhaps we'll hold that what she ought to believe about when the party starts depends on when Peter's parties have started in the past (which we take her to remember) or on what Peter's email said (which again, we take her to remember), or just on her best guess as to what Peter said (and we'll take her to be in a position to know what her best guess is), or on some combination of the above sorts of facts.[47] Whatever the new local supervenience base is for facts about what Justine ought to believe, as long as Justine counts as in a position to know the truth about facts in the new base, there needn't be any obstacle to her being in a position to know what she has justification to believe.

Here's the more general picture into which the above example fits. Each context will determine a local supervenience base for facts about what a subject ought to believe—a set of propositions B such that, within the local set of worlds taken seriously in the context, there are no differences in what the subject ought to believe without differences in which propositions in B are true. Moreover, it is a constraint on such bases that if B is a supervenience base for S's justification in context C, then S counts as in a position to know which are the truths in B, relative to context C. This way, we will never say that the facts that determine what a subject ought to believe are facts that the subject is not in a position to know. We can think of B as a sort of cognitive home—a realm of propositions such that the subject is always in a position to know which are the truths in that realm.

But because we do not believe in cognitive homes of the traditional sort, we'll admit that whatever B is, we *will* be able to raise reasonable doubts about S's ability to know which are the truths in B. To the extent that we take such doubts seriously, the context will shift, and in our new context C', B will not be the local supervenience base for facts about what S has justification to believe. Rather, there will be some new supervenience

[47] In any of these cases, we might take it that she oughtn't have outright beliefs about the start time of the party, but should instead distribute her confidence over various possibilities. This is compatible with (Access), insofar as we take Justine to be in a position to know how she should distribute her confidence.

base B', such that S *does* count as in a position to know which are the truths in B'. In each context, we'll count S as having *some* cognitive home (i.e., there will be *some* supervenience base for facts about S's justification, such that S is in a position to know which are the truths in the base), but *which* class of facts constitutes S's home will vary from context to context; because S's cognitive home is apt to shift its location whenever doubts of certain sorts are taken seriously, we might call it a mobile home.

I've sketched a picture for how the contextualist (or relativist, or expressivist) might evade *one* sort of argument against a version of access internalism—namely, she can show that it doesn't depend on a very strong version of traditional foundationalism according to which there is some distinctive class of propositions (e.g., concerning our mental lives) that form a supervenience base for facts about epistemic justification, and whose truth values we are always (in all contexts) in a position to know. But while opponents of this brand of internalism often do stress the (apparent) connection between it, and the traditional foundationalist view just mentioned,[48] it's certainly not the *only* argument frequently offered against (Access), nor is it the most influential—that distinction belongs to Timothy Williamson's anti-luminosity argument (2000, ch. 4).

As I've already mentioned, (Access) is a luminosity claim—it says that being such that you ought to believe that P is a luminous condition, for all P. And Williamson's argument purports to show that there are no non-trivial luminous conditions. There's been a great deal written on the anti-luminosity argument, and it would be beyond the scope of this paper to get into the literature on that argument in any great depth.[49] But I do want to show how contextualism suggests a natural place for friends of (Access) to look in responding to the anti-luminosity argument.

Williamson's argument turns on a margin-for-error principle, to the effect that if one is to know that P in case C, one must not falsely believe that P in any case "close" to C.[50] To see how contextualism suggests a response, it may help to see Williamson's position as a version of the relevant alternatives approach to knowledge.[51] On relevant alternatives approaches, it is a necessary condition for knowing that P that one avoid falsely believing that P within some range of relevant possibilities. On Williamson's view in particular, the range of relevant cases itself varies from case to case (in any case C, the relevant cases are the cases close to C), but not from context to context. This leads to anti-luminosity via a *reductio*. Because the range of relevant cases varies from case to case, the range of alternatives that are relevant to cases that are relevant to a given case

[48]See Williamson (2000), Srinivasan (Ms.).

[49]For some criticisms of the argument see, Neta and Rohrbaugh (2004), Weatherson (2004) Berker (2008), Cohen (2010), Vogel (2010). For a recent defense of luminosity against (many of) these criticisms, see Srinivasan (2013).

[50]I put "close" in scare quotes to highlight that Williamson does not presuppose that the sense of closeness relevant to the margin-for-error principles he uses is an familiar one that is used in judgments of closeness outside the context of epistemology. While this makes it easier to avoid putative counterexamples to margin-for-error principles and the closely related safety principle—faced with a putative example in which somebody knows that P in some case C, even though they falsely believe that P in a case C′ that seems to be close to C, one can deny that C′ is close in the sense relevant to knowledge—it also makes it harder to motivate the principles. This criticism is stressed by Cohen (2010).

[51]The approach is originally due to Dretske (1970), but it has been discussed by many authors since.

C is strictly larger than the range of cases that are relevant to C. But if some condition ϕ is luminous, then (via the arguments provided by Williamson), if one is in ϕ in case C, one must also be in ϕ in all cases relevant to cases relevant to cases relevant to cases ...(*ad infinitum*) relevant to C. Because each iteration of "cases relevant to" widens the range of possibilities, only trivial conditions—conditions that hold everywhere or nowhere in logical space—have a hope of being luminous.

The natural contextualist response is to hold that the range of relevant cases varies from context to context, but within a context, it does not systematically vary from case to case.[52] Roughly, to know that P (by the standards of context α) in case C, one must avoid falsely believing that P in all relevant (by the standards of α) cases. But if each context fixes a single set of relevant cases (that does not vary from case to case), then iterating "cases relevant to" doesn't widen the set of relevant alternatives. For some condition to hold in all the cases that are relevant to cases relevant to cases relevant to ...(*ad infitium*) some case C (by the standards of a given context α) is no harder than for it to hold in all the cases relevant (by the standards of α) to C, since relevance is relativized not to a case, but instead to a context. There are extant contextualist views with essentially this structure.Lewis (1996) defends a view on which knowing that P requires being in a state that—within the range of not properly ignored (i.e., relevant) possibilities—one is only in if P holds, and where the range of relevant possibilities varies from context to context, but not systematically from case to case. [53]

Is the resulting view defensible? After all, it will have to accept that there are counterexamples to margin-for-error principles, and won't this be counterintuitive? Yes, but the strategies I've discussed earlier in this essay suggest a natural reply. Suppose we try to find an intuitively plausible counterexample to Williamson's margin-for-error principle. This would require identifying a case C, a subject S, and a proposition P such that it seems to us as if S knows that P in C, even though there is a nearby and relevantly similar case C$'$ which S falsely believes that P. But if explicitly considering a case C$'$ tends to make that case salient, and thus induces a context shift so that C$'$ counts as a relevant possibility, then we shouldn't expect to be able to find such a counterexample. Once we explicitly consider the possibility C$'$, that possibility will count as relevant. Because, in C, S is not in a state that allows her to rule out the relevantly similar possibility that she is in C$'$, in the new context in which C$'$ counts as a relevant possibility, S will not count as knowing that P in case C.

Counterexamples to Williamsonian margin-for-error principles, then, might be "elusive," in the sense described by Lewis (1996). While in any given context, there may

[52]In particular, it's not as if each iteration of "relevant to" brings us to a wider set of possibilities, as it must if relevance is closeness.

[53]A number of recent defenders of iteration principles in epistemology (e.g., the KK principle) have made versions of this point. Stalnaker (2009), Cohen and Comesaña (2013), and Greco (Forthcoming) each argue that the intuitive motivation for margin-for-error principles—that one can't know if one could've easily been wrong—can be captured within formal systems that do not vindicate the assumptions required by the anti-luminosity argument. Greco in particular pursues a strategy similar to the one urged here, in which contextualism plays a key role in allowing us to retain the intuitive motivation for margin for error principles, while rejecting the versions of them that Williamson appeals to in the anti-luminosity argument.

be instances of the principles that are false, explicitly considering those instances will typically bring us into a new context in which they are true. So it is unsurprising if the principles seem true—their instances will typically be such that, once they are explicitly considered, they *are* true. This is effectively the same strategy I suggested on behalf of the Bayesian account of defeat, and that Raffman and Fara suggest in defense of the idea that vague predicates determine precise cutoffs (in given contexts of use). In each case, some principle seems intuitive, and its intuitiveness is explained away by holding that while counterexamples exist, there are systematic reasons why considering them changes the context in a way such that they no longer count as counterexamples.

While I've tried to show how the contextualist *can* defend (Access), I certainly haven't shown that she *should*; I don't take the availability of the strategies I've presented in this section to come close to constituting a knock-down argument in favor of (Access). I do take it, however, to change the dialectic concerning (Access) in an important way. The truth of (Access) can *seem* to require the truth of a very strong version of traditional foundationalism. While some friends of (Access) have accepted this and endorsed such strong versions of traditional foundationalism, most philosophers will regard this as too big a bullet to bite. Once we see that (Access) does not depend on such a strong version of traditional foundationalism, but can get by appealing only to a more modest contextualist foundationalism, we can evaluate the plausibility of (Access) independently of the plausibility of strong versions of traditional foundationalism. Moreover, (Access) is incompatible with the anti-luminosity argument. But again, to the extent that we have some sympathy for "elusiveness" moves in other cases (knowledge more generally, and precise cutoffs for vague predicates), such moves suggest a natural reply to the anti-luminosity argument. So the reasons to reject (Access) aren't as decisive as is typically thought. Still, while I adduced some *prima facie* motivations for (Access) at the beginning of this section, I haven't tried to show how a full accounting of its costs and benefits would turn out.

4 Conclusion

The idea that contextualism can provide a sort of middle way in the debate between coherentism on the one hand, and traditional foundationalism on the other, is not new. But with the exception of some discussion of traditional skeptical problems, there has been relatively little work done exploring how *other* epistemological debates look once we take this middle way.[54] I regard this as an oversight. Contextualist foundationalism promises to illuminate debates about Bayesianism and defeat, debates about internalism in epistemology, and probably many others as well.

[54]Much of Michael Williams' work involves using a sort of contextualist foundationalism to argue that global skeptical challenges are somehow illegitimate. He has not, to my knowledge, discussed the relationship between contextualist foundationalism and Bayesianism, or attempted to use contextualist foundationalism to defend anything like (Access)—in fact, he endorses a qualified sort of externalism in Williams (Forthcoming). Ram Neta's work (2003; 2004; 2005) also involves something like contextualism about which beliefs are foundational, but to my knowledge he has not applied his version of contextualism to the topics I've addressed in the last two sections of this paper.

Bibliography

Aikhenvald, Alexandra. 2004. *Evidentiality*. Oxford University Press.

Alston, William P. 1976. Has Foundationalism Been Refuted? *Philosophical Studies* 29:295.

Annis, David B. 1978. A Contextualist Theory of Epistemic Justification. *American Philosophical Quarterly* 15:213–219.

Berker, Selim. 2008. Luminosity Regained. *Philosophers' Imprint* 8:1–22.

Blome-Tillmann, Michael. 2009. Knowledge and Presuppositions. *Mind* 118:241–294.

Carnap, Rudolf. 1967. *The Logical Structure of the World [and] Pseudoproblems in Philosophy*. London, Routledge K. Paul.

Chrisman, Matthew. 2007. From Epistemic Contextualism to Epistemic Expressivism. *Philosophical Studies* 135:225–254.

Christensen, David. 1992. Confirmational Holism and Bayesian Epistemology. *Philosophy of Science* 59:540–557.

Cohen, Stewart. 1987. Knowledge, Context, and Social Standards. *Synthese* 73:3–26.

———. 1988. How to be a Fallibilist. *Philosophical Perspectives* 2:91–123.

———. 2010. Luminosity, Reliability, and the Sorites. *Philosophy and Phenomenological Research* 81:718–730.

Cohen, Stuart, and Juan Comesaña. 2013. Williamson on Gettier Cases and Epistemic Logic. *Inquiry* 56:15–29.

Conee, Earl, and Richard Feldman. 2001. Internalism Defended. In Hilary Kornblith, editor, *Epistemology: Internalism and Externalism*. Oxford: Blackwell. Reprinted in Conee and Feldman 2004.

———. 2004. *Evidentialism*. Oxford: Oxford University Press.

DeRose, Keith. 1995. Solving the Skeptical Problem. *The Philosophical Review* 104:1–52.

———. 2009. *The Case for Contextualism*. Oxford University Press.

Dretske, Fred I. 1970. Epistemic Operators. *Journal of Philosophy* 67:1007–1023.

Fantl, Jeremy, and Matthew McGrath. 2009. *Knowledge in an Uncertain World*. Oxford: Oxford University Press.

Fara, Delia Graff. 2000. Shifting Sands: An Interest Relative Theory of Vagueness. *Philosophical Topics* 28:45–81.

Feldman, Richard. 2005. Respecting the Evidence. *Philosophical Perspectives* 19(1):95–119.

Field, Hartry. 2009. Epistemology Without Metaphysics. *Philosophical Studies* 143.

Finlay, Stephen, and Gunnar Björnsson. 2010. Metaethical Contextualism Defended. *Ethics* 121:7–36.

Garber, Daniel. 1982. Old Evidence and Logical Omniscience in Bayesian Confirmation Theory. In John Earman, editor, *Testing Scientific Theories*. University of Minnesota Press.

Gibbard, Allan. 2005. Truth and Correct Belief. *Philosophical Issues* 15:338–350.

Goodman, Nelson. 1983. *Fact, Fiction, and Forecast*. Harvard University Press.

Greco, Daniel. 2014a. Could KK Be OK? *Journal of Philosophy* 111:169–197.

———. 2014b. Iteration and Fragmentation. *Philosophy and Phenomenological Research* 88.

Harman, Gilbert. 2001. General Foundations Versus Rational Insight. *Philosophy and Phenomenological Research* 63:657–663.

Hedden, Brian. 2012. Options and the Subjective Ought. *Philosophical Studies* 158:343–360.

Horowitz, Sophie. Forthcoming. Epistemic Akrasia. *Noûs* .

Ichikawa, Jonathan Jenkins. 2013. Basic Knowledge and Contextualist "E = K". *Thought: A Journal of Philosophy* 2:282–292.

Kolodny, Niko, and John MacFarlane. 2010. Ifs and Oughts. *Journal of Philosophy* 107:115–143.

Kotzen, Matthew. Unpublished Manuscript. A Formal Account of Evidential Defeat .

Lewis, David. 1996. Elusive Knowledge. *Australasian Journal of Philosophy* 74:549–67. Reprinted in Lewis 1999.

———. 1999. *Papers in Metaphysics and Epistemology*. Cambridge: Cambridge University Press.

MacFarlane, John. 2011a. Epistemic Modals Are Assessment-Sensitive. In Andy Egan and Brian Weatherson, editors, *Epistemic Modality*. Oxford University Press.

———. 2011b. Relativism and Knowledge Attributions. In *Routledge Companion to Epistemology*. Routledge.

Mayo, Deborah. 1996. *Error and the Growth of Experimental Knowledge*. University of Chicago Press.

Neta, Ram. 2003. Contextualism and the Problem of the External World. *Philosophy and Phenomenological Research* 66:1–8211.

———. 2004. Perceptual Evidence and the New Dogmatism. *Philosophical Studies* 119:199–214.

———. 2005. A Contextualist Solution to the Problem of Easy Knowledge. *Grazer Philosophische Studien* 69:183–206.

Neta, Ram, and Guy Rohrbaugh. 2004. Luminosity and the Safety of Knowledge. *Pacific Philosophical Quarterly* 85:396–406.

Plunkett, David, and Tim Sundell. Forthcoming. Disagreement and the Semantics of Normative and Evaluative Terms. *Philosophers' Imprint* .

Pryor, James. Forthcoming. Problems for Credulism. In *Seemings and Justification: New Essays on Dogmatism and Phenomenal Conservatism*. Oxford University Press.

Quine, Willard Van Orman. 1951/1953. Main Trends in Recent Philosophy: Two Dogmas of Empiricism. *The Philosophical Review* 60:20–43. Reprinted, with alterations, in Quine 1953/1980.

———. 1953/1980. *From a Logical Point of View*. Cambridge, MA: Harvard University Press, revised second edition.

Raffman, Diana. 1994. Vagueness Without Paradox. *Philosophical Review* 103:41–74.

———. 1996. Vagueness and Context-Relativity. *Philosophical Studies* 81:175–192.

Rayo, Agustín. Forthcoming. A Plea for Semantic Localism. *Noûs* .

Schaffer, Jonathan. 2004. From Contextualism to Contrastivism. *Philosophical Studies* 119:73–104.

———. 2007. Knowing the Answer. *Philosophy and Phenomenological Research* 75:383–8211.

Schwitzgebel, Eric. 2011. *Perplexities of Consciousness*. MIT Press.

Sepielli, Andrew. Forthcoming. Subjective Normativity and Action Guidance. *Oxford Studies in Normative Ethics* .

Smithies, Declan. 2012. Moore's Paradox and the Accessibility of Justification. *Philosophy and Phenomenological Research* 85:273–300.

Sosa, David. 2009. Dubious Assertions. *Philosophical Studies* 146:269–272.

Srinivasan, Amia. 2013. Are We Luminous? *Philosophy and Phenomenological Research* 87.

———. Unpublished Manuscript. What's in a Norm? .

Stalnaker, Robert. 2008. *Our Knowledge of the Internal World*. Oxford: Oxford University Press.

———. 2009. On Hawthorne and Magidor on Assertion, Context, and Epistemic Accessibility. *Mind* 118:399–409.

Stanley, Jason. 2005. *Knowledge and Practical Interests*. Oxford University Press.

Vogel, Jonathan. 2010. Luminosity and Indiscriminability. *Philosophical Perspectives* 24:547–572.

Weatherson, Brian. 2004. Luminous Margins. *Australasian Journal of Philosophy* 82:373 – 383.

Weisberg, Jonathan. 2009. Commutativity or Holism? A Dilemma for Conditionalizers. *British Journal for the Philosophy of Science* 60:793–812.

———. Forthcoming. Updating, Undermining, and Independence. *British Journal for the Philosophy of Science* .

Williams, Michael. 1999—First edition published in 1977. *Groundless Belief*. Princeton University Press.

———. Forthcoming. Skepticism, Evidence and Entitlement. *Philosophy and Phenomenological Research* .

Williamson, Timothy. 1994. *Vagueness*. London: Routledge.

———. 2000. *Knowledge and its Limits*. Oxford University Press.

———. 2005. Contextualism, Subject-Sensitive Invariantism and Knowledge of Knowledge. *Philosophical Quarterly* 55:213–235.

Wittgenstein, Ludwig. 1969. *On Certainty*. Basil Blackwell, Oxford.