Mark Schroeder
University of Southern California
October 19, 2008

what does it take to 'have' a reason?
forthcoming in reisner and steglich-peterson, eds., *Reasons for Belief*

> If I believe, for no good reason, that P and I infer (correctly) from this that Q, I don't think we want to say that I 'have' P as evidence for Q. Only things that I believe (or could believe) rationally, or perhaps, with justification, count as part of the evidence that I have. It seems to me that this is a good reason to include an epistemic acceptability constraint on evidence possessed…[1]

It is a truism that adopting an unjustified belief does not put you in a better evidential position with respect to believing its consequences. This truism has led many philosophers to assume that there must, at a minimum, be a justification condition (and perhaps even a knowledge condition) on what it takes to count as having evidence. This is the best (or only) possible explanation of the truism, these philosophers have believed. This paper explores an alternative explanation for the truism. According to the alternative explanation that I will offer, unjustified beliefs do not put you in a better evidential position with respect to believing their consequences because any evidence you have in virtue of having an unjustified belief is guaranteed to be defeated. Since the lack of justification for a belief guarantees its defeat, I will suggest, we don't need to postulate a special justification condition (much less a knowledge condition) on what it takes to count as having evidence.

Why is this important? It is important because the assumption that there must be a justification condition (or perhaps a knowledge condition) on what it takes to count as having evidence places a *high bar* on what it takes to have evidence - such a high bar that it is difficult to see how this bar could be met in the case of basic, perceptually justified beliefs. As a result, the high bar set by this condition plays a fundamental role, I will claim, in central features of a core dialectic from the epistemology of basic perceptual belief which plays a central role in the debates between internalism and externalism, foundationalism and coherentism, and rationalism and empiricism. Given the degree to which this dialectic is shaped by the assumption of the high bar on what it takes to have evidence, I believe that it is well worth

---

[1] Feldman [1988, 227].

taking seriously what an alternative view might look like. My aim in this paper is not to defend this alternative picture or its virtues in detail; it is simply to make room for the alternative.

I   clarifications: having evidence

If Max is smiling, that is evidence that he is happy. So if you can see Max, and he visually appears to you to be smiling, then you have some evidence that he is happy. Similarly, if you find out by testimony that Max is smiling, you have some evidence that he is happy. But if Max is a thousand miles away from you and you have no access to whether he is smiling or not, then even though the fact that he is smiling is evidence that he is happy, it is not evidence that you have. The difference between evidence that you have and evidence that you don't have is crucially important. For example, only evidence that you have affects what it is rational for you to believe. If you can see Max, and he visually appears to you to be smiling, that affects how rational it would be for you to believe that Max is happy. But if Max is thousands of miles from you and you have no access to whether he is smiling or not, the fact that he is in fact smiling makes no difference whatsoever to how rational it would be for you to believe that he is happy. Similarly and also importantly, your beliefs can be *based* only on evidence that you have; if Max is thousands of miles away and you have no access to whether he is smiling or not, the fact that he is smiling can't be the evidence *for which* or *on the basis of which* you believe that he is happy. Whereas if you can see Max, and he visually appears to you to be smiling, it then becomes possible for this evidence to be that for which you believe that he is happy. Because only evidence that you have affects what it is rational for you to believe, and because you can believe on the basis only of evidence that you have, the notion of evidence that you have is very important in epistemology.

Talk about having evidence that you have makes it sound like when you have some evidence, what happens is that there is something which is, independently, *evidence*, and which, moreover, you *have* - which is in your possession or grasp in some way. This idea is what I call the Factoring Account, because it factors 'having evidence' into two parts – 'having', and 'evidence'. I don't actually think that the Factoring Account is true (I have previously published a paper explaining why not – see Schroeder [2008]). To begin to get suspicious about whether the Factoring Account should be obvious or not, compare talk about having evidence to talk about having fathers. Todd has a father; his father is Larry. But Larry is not, independently, a father, of whom Todd happens to be in possession. What makes Larry the father Todd has is simply that he stands to Todd in the *father-of* relation. So certainly not all talk about 'having' something is talk about possession. Moreover, the Factoring Account has a strong consequence: if Max is

not really smiling, then it's not really evidence that he is happy, that he is smiling. That is, talk about what is evidence *simpliciter* is naturally understood as *factive*. So if you can have evidence only if it is really evidence, then you never have evidence, unless what you believe or perceive to be the case is true. This is a big conclusion to pull out of such a little hat.

In opposition to the Factoring Account is the Two Relations Account, which I favor instead. According to the Two Relations Account, there are simply two relations that we need to be careful about – the *objective* evidence relation, and the *subjective* evidence relation. When I started off by noting that if Max is smiling, that is evidence that he is happy, what I said was true about evidence in the objective sense, and when I spoke about what evidence you *have* that Max is happy, what I said was true about evidence in the subjective sense. Objective evidence must be true, but subjective evidence need not be true; subjective evidence must be in some sense yet to be elucidated in the subjects 'grasp', but objective evidence need not be in the subject's grasp. Moreover, according to the Two Relations Account, these two relations are not unconnected; on the contrary, they are related, very roughly, by the twin facts that any objective evidence that is in the subject's grasp (in the appropriate sense) is subjective evidence, and that any subjective evidence that is true is objective evidence.

Unlike the Factoring Account, which treats subjective evidence as a *restriction* on objective evidence, the Two Relations Account leaves open that subjective evidence need not be true. But both accounts appeal at a crucial juncture to a kind of grasp or possession relation. The Factoring Account appeals to this possession relation when it claims that subjective evidence is (objective) evidence that you *have*. The Two Relations Account appeals to this possession relation when it claims that subjective evidence must be in the subject's 'grasp'. It is the nature of this grasping or possession relation that I will be interested in, in this paper. The truism that adopting an unjustified belief does not put you in a better evidential position with respect to believing its consequences leads most philosophers to conclude that bare belief is insufficient for the *having* relation involved in subjective evidence. These philosophers adopt the view that, at least with respect to inferentially justified beliefs (beliefs justified on the basis of other beliefs), your evidence is restricted to what you *justifiably* believe, or even to what you *know*. It is this conclusion that I want to illustrate how to resist, in this paper. What I will be trying to show, is that since any evidence you have in virtue of your unjustified beliefs is guaranteed to be defeated, it does not matter (at least for the purposes of accommodating the truism) whether we count their contents as part of what you 'have' in the relevant sense, or not. Consequently we can take whichever view leads to the best overall theory.

2   the setup: two arguments for a high bar, and two readings of the truism

As I indicated at the beginning, many philosophers assume that we need to place a substantial restriction or 'high bar' on this *having* relation. In order to get a sense for how this high bar works, we can introduce some simple notation for thinking about different views about the having relation as ranging from very liberal to very restrictive, as in the diagram, below. For ease of notation, I will use 'H=K' as an abbreviation for the thesis that the having relation is knowledge.[2] Similarly, I will use 'H=JB' as an abbreviation for the thesis that the having relation is justified belief, and 'H=B' will denote the thesis that the having relation is simply belief. The thesis for which I will be trying to make room, in this paper, if not to argue for outright, is H=Pres – that the having relation is that of having a *presentational* attitude toward a proposition. (Presentational attitudes, as I will understand the term, are attitudes which present their content to their subject as being true, which I understand to include both belief and perceptual experience (both veridical and non-veridical), but not desire, wonder, supposition, assumption, or many other propositional attitudes.)

$$\text{H=Pres} \quad \text{H=B} \quad \Big| \quad \text{H=JB} \quad \text{H=K}$$
$$\longleftarrow \quad \text{liberal} \qquad \qquad \text{restrictive} \quad \longrightarrow$$

There are two ways of understanding the motivation for postulating the high bar on the having relation: one direct, and one indirect. The indirect argument begins with the observation of our truism: that adopting an unjustified belief does not put you in a better evidential position with respect to believing its consequences. It then argues, by way of some sort of auxiliary premise, such as the premise that you have evidence for *p* by being in a state of mind *S* only if being in *S* puts you in a better evidential position with respect to believing *p*, that since unjustified beliefs do not put you in a better evidential position with respect to believing their consequences, they are not ways of having evidence for their consequences (similar points go for propositions which would be supported, but are not outright consequences, of the contents of these beliefs – a complication which I'll henceforward ignore, just to simplify things).

This argument has one impeccable premise – the truism with which we began. But it is indirect, because it relies on an auxiliary premise connecting having evidence with improving one's evidential position. But there is also a direct argument for the high bar on the having relation. It proceeds by pumping the direct intuition that someone who has no other reason to believe *p* does not have any evidence

---

[2] Williamson [2000] uses 'E=K' as his name for a quite similar thesis.

that *p*, simply in virtue of unjustifiedly believing something which has *p* as a consequence (the same points about generalizing apply as before, as well as the note that as in this case, I will sometimes be shamelessly mixing use with mention in order to avoid awkward locutions). This argument skips the auxiliary premise, but the intuition on which it relies is correspondingly more controversial. In section 3 I'll respond to the indirect argument by offering an alternative explanation for the truism – an explanation which incidentally sheds some light on why the auxiliary premise is false. Then in section 4 I'll take up the direct argument, and argue, as I have elsewhere, that negative intuitions about reasons or evidence, like those it relies on, are systematically misleading under the same circumstances which make the indirect argument's auxiliary premise false – and in fact, in a wider range of circumstances than these. The direct argument, I'll therefore be arguing, is not independent from the indirect argument at all, and my earlier response to the indirect argument therefore suffices to set it aside as well.

Before I get to either of these arguments, however, we require yet one more important clarification as to what, exactly, is a truism, about our truism. Epistemologists importantly distinguish between *having* a justification to believe that *p*, and being justified *in* believing that *p*. This is sometimes called the distinction between *propositional* and *doxastic* justification. The idea is that *p* is propositionally justified for Maria when she *has* a justification to believe it, even if she does not believe it, or believes it, but not for the reasons involved in that justification at all, but instead (coincidentally) believes it because she read it in her horoscope. (Assume for the sake of argument that *p* is among the propositions for which Maria's horoscope does not provide evidence.) Whereas *p* is doxastically justified for Maria only if Maria believes that *p*, and does so on the basis of her justification to believe that *p*.

Given this important distinction, a belief that *p* can be unjustified either in the sense that it lacks propositional justification, or in the sense that the subject does not believe it on the basis of that justification. And that means that if we assume that a 'better evidential position' reflects some improvement in one of our two senses of 'justification', our truism is four ways ambiguous:

PP   When *p* lacks propositional justification, it does not contribute to the agent's propositional justification for its consequences.

PD   When *p* lacks propositional justification, it does not contribute to the agent's doxastic justification for its consequences.

DP   When *p* lacks doxastic justification, it does not contribute to the agent's propositional justification for its consequences.

DD   When *p* lacks doxastic justification, it does not contribute to the agent's doxastic justification for its consequences.

5

Of these four, I believe that only three are truisms. If Maria has an adequate propositional justification to believe that p, but her belief that p lacks doxastic justification only because she believes it for the wrong reasons, I believe that she can still be propositionally justified in believing p's consequences. Her propositional justification for believing that p will also justify her in believing p's consequences, on this view, though in accordance with DD, she will not be doxastically justified in believing those consequences on the basis of p unless she believes p for the right reasons.

Moreover, though I take all three of PP, PD, and DD to be truisms, I think that of these, PP is the most basic and most important. DD can be explained on the basis of the assumption that to be doxastically justified, there must be no gaps in the doxastic justification of the premises on the basis of which it is held. This is an assumption specifically about doxastic justification, and needn't have consequences for the notion of having evidence – in fact, if we held that DD is true because when p lacks doxastic justification, the agent does not have p as evidence for q, we would intuitively predict that DP should also be true – which, as I have just suggested, it is not. Similarly, since doxastic justification implies propositional justification, PP's truth will suffice to explain the truth of PD. Consequently, in section 3 I am going to focus on an argument which seeks to predict and explain the truth of PP – the most basic and important of the readings of our truism. It is this reading which I believe plays the central role in the indirect argument for a high bar on the having relation, so it is this reading which I will try to explain.

3   the alternative explanation: lack of justification guarantees defeat

In this section I offer my alternative explanation of PP – the truism that when a belief lacks propositional justification, it does not contribute to the propositional justification for its consequences. Actually, I'll argue for something stronger; what I'll be arguing is for the schema *Lack of Justification Guarantees Defeat*:

> LJGD   When an agent's belief that p lacks propositional justification, it does not contribute to her propositional justification to believe that q.

Since this schema holds no matter which sentences we substitute for 'p' and 'q', it is not just that unjustified beliefs do not contribute to the propositional justification for their consequences; they do not contribute to the justification for anything at all.

I give this schema its name – *Lack of Justification Guarantees Defeat* – in order to highlight a feature of its explanatory structure. Rather than explaining our truism by restricting the having relation, my proposal is that we explain the truism by appeal to general predictions about evidence being *defeated*. Epistemologists

6

importantly distinguish between two kinds of 'defeat'; a piece of evidence has a *countervailing* defeater when there is better contrary evidence – for example, though the fact that Max told you that he would come is evidence that he would come, the fact that you can see with your own eyes that he did not come is even better evidence to the contrary. In this example, the fact that you can see with your own eyes that Max did not come is a countervailing defeater for the evidence that he would come. A piece of evidence thus suffers from countervailing defeat when it 'loses' to better evidence on the other side.

In contrast, a piece of evidence has an *undercutting* defeater when there some further consideration which mitigates its force – perhaps so much so, that it does not carry any weight at all. For example, the fact that Max is often deceitful is an undercutting defeater for the fact that he said that he would come, as evidence that he would come. If Max is often deceitful, then what he says is simply less good evidence about what is true, than it otherwise would be. A piece of evidence can be undercut without being countervailed; Max's word may still be all that you have to go by, and if he is often *not* deceitful, then it may still count for something. But when it is undercut, it counts for less, and maybe for nothing at all. (In evaluating whether a piece of evidence is undercut or countervailed, keep in mind that sometimes one and the same consideration does *both* – for example, if Max is a reliable liar, then not only does what he says not count in favor of its truth, what he says, together with the fact that he is a liar, counts *against* its truth. In such a case there is an interaction between an undercutting and countervailing evidence.)

Ordinary cases of undercutting defeaters are cases of what I call *specific* defeat. In these cases, *d* defeats *p* as a reason to believe *q*, but does not defeat *p* as a reason to believe other things. For example, that Tweety is a bird is evidence that Tweety flies. It is also evidence that Tweety has feathers. That Tweety nests in Antarctica is an undercutting defeater for the former piece of evidence, but not for the latter. It defeats the evidence provided by the proposition that Tweety is a bird for the conclusion that Tweety flies, but it does not defeat the evidence provided by the proposition that Tweety is a bird for the conclusion that Tweety has feathers. Such defeaters I'll call *specific*, because their defeat is specific to particular conclusions that we might draw on the basis of some particular piece of evidence.

In contrast, the kind of defeat postulated by LJGD is what I'll call *general* defeat. General defeaters don't just defeat one or another particular conclusion that we might draw on the basis of a piece of evidence; they defeat any such conclusion. A picture may help to clarify the distinction:

In the picture, the 'open' ended arrows ($\longrightarrow$) represent the direction of reasons to believe P, to believe Q, and to believe R, and the 'closed' ended arrows ($\longrightarrow$) represent different kinds of defeater. As illustrated, defeaters may be *countervailing* or *undercutting*. Countervailing defeaters attack the conclusion directly; they are reasons to *not* believe something – usually, reasons to believe its negation.[3] Undercutting defeaters, in contrast, do not attack the conclusion directly, but rather attack the reasons which support it. Specific undercutting defeaters attack these reasons 'downstream' from the proposition which constitutes the reason, and as a result their effects are specific to particular conclusions. Whereas general defeaters, in contrast, attack these reasons 'upstream' from the proposition which constitutes the reason, which means that they may undermine everything that P is a reason for.

You might think that general defeaters are too different from specific defeaters to really count as a kind of defeat. You might think that their very generality makes it more natural to describe considerations which have general defeaters as not evidence at all, rather than as evidence that has been defeated. I'm not so much concerned about how it is natural to describe them, however, as with the underlying explanation of the phenomenon. After offering the following argument for the thesis that beliefs which lack propositional justification *are*, in fact, guaranteed to have general defeaters, I'll therefore go back and explain why I think this argument effectively explains these general defeaters by the same general principles that suffice to explain how specific defeaters work, as well.

Finally, we get to the argument for LJGD. It is quite simple:

P1   When an agent's belief that *p* lacks propositional justification, she has insufficient
       reason to believe that *p*.
P2   When an agent has insufficient reason to believe that *p*, she has conclusive reason
       to not believe that *p*.
P3   When an agent has conclusive reason not to believe that *p*, she has conclusive
       reason not to take it into account that *p* in her reasoning about whether *q*.

---

[3] Although see Schroeder [unpublished]; some reasons not to believe that *p* might be reasons to withhold, rather than reasons to believe that ~*p*.

P4   When an agent has conclusive reason not to take it into account that *p* in her reasoning about whether *q*, it is irrational for her to take it into account that *p* in her reasoning about whether *q*.

P5   An agent's belief that *p* contributes to her propositional justification to believe that *q* only if it is rational for her to take it into account that *p* in her reasoning about whether *q*.

LJGD   When an agent's belief that *p* lacks propositional justification, it does not contribute to her propositional justification to believe that *q*.

This argument has five premises. Of these, I take P1 to be a highly compelling principle about propositional justification and P2 to follow from the definitions of 'sufficient' and 'conclusive'. I take P3 to be independently highly compelling; we should not rely in our reasoning on things that we don't believe to be the case. P4 I hope is close to a platitude, so long as we keep in mind that we are throughout talking about *subjective* reasons – reasons the agent 'has' – as opposed to objective reasons, which might be conclusive even though the agent is unaware of them and could be rational in flouting. P4 is certainly consistent with the idea that there is much more to rationality than complying with conclusive subjective reasons. And finally, P5 looks to me like it is at least highly plausible – certainly the normal way in which beliefs contribute to the justification for other beliefs is because it is rational to take their contents into account in reasoning. If the belief that *p* does not contribute to the justification for believing *q* in this way, then it is not clear how it does contribute to it.

So together, P1-P5 provide an argument from highly plausible premises for LJGD, interpreted as the thesis that lack of propositional justification guarantees defeat. This argument shows how we can rule out unjustified beliefs helping to justify other beliefs, without the need to postulate a special justification condition on the 'having' relation involved in having evidence. According to this argument, *even if* you can 'have' evidence by having a propositionally unjustified belief, it is guaranteed to be defeated anyway.

It is natural to question whether the argument that I have just provided gives us an independent *explanation* of our truism, or just predicts that our truism is, in fact, true – a fact which we should explain by positing a high bar on the having relation. I believe that the argument I have just given does, in fact, provide an independent explanation of the truism, because I believe that it illustrates that the very same general principles which explain the behavior of undercutting defeaters in general, predict and explain why there would be general undercutting defeaters in these cases. This is because in general, undercutting defeaters work by making it rational to take the proposition which they defeat less into account in reasoning toward their conclusion. To see how this works, let's go back to one of our examples. The fact

9

that Tweety lives in Antarctica is an undercutting defeater for the fact that Tweety is a bird, as evidence that Tweety flies. If you are aware that Tweety lives in Antarctica, this also makes it rational (given other common background knowledge) for you to place very little weight on the fact that Tweety is a bird, in reasoning about whether Tweety flies.

The structure of the argument that I have just given aims to show that the very same thing is going on when you have an unjustified belief – namely, that you have some reason which makes it rational to place very little – even no – weight on the content of your unjustified belief in your reasoning. Which reason is this? It is the sufficient reason not to believe the proposition in the first place – which in most cases is largely constituted by the evidence that you have against it, which made it an unjustified thing to believe in the first place.[4] So the idea of the argument is that the very same principles which explain how specific undercutting defeaters work – by making it rational to pay the evidence they undercut less attention in reasoning – also serve to explain how general undercutting defeaters work – by making it rational to pay the evidence they undercut less attention in reasoning. Consequently, the idea is that our truism is simply a consequence of these general features of the behavior of undercutting defeaters. It is therefore not something of which we need to provide a special explanation, by postulating special high bars on what it takes to have evidence.

4   negative intuitions about reasons/evidence and the direct argument

So far, I take myself to have responded to the indirect argument for the high bar on having evidence. I now turn to the direct argument, which trades on eliciting the direct intuition that someone (who has no other evidence that q) who believes that p only unjustifiedly *has no* reason to believe that q – or, put in terms of evidence, that she *has no* evidence that q. My response constitutes an error theory about negative intuitions about reasons and evidence under certain circumstances which are present in this case. In particular, I'll show that when reasons or evidence are undercut or otherwise very low in weight – at least on contrast to the competing reasons on the other side – it is quite easy to elicit the intuition that someone does not have a reason to do something, even when she really does have such a reason, albeit one of low weight. Here I'll make the argument directly in terms of intuitions about reasons, rather than in terms of intuitions about evidence, though I believe that nearly identical considerations apply. (Readers who can anticipate how this argument will go or who are familiar with versions of the same point which I have presented elsewhere may

---

[4] Although again see Schroeder [unpublished] for a discussion of how some reasons not to believe that p might not be evidence for $\sim p$, but rather reasons to withhold.

profitably skip directly to section 5.)

The reason why the direct argument does not do any more work than the indirect argument, is that our intuitive judgments about what there is no reason to do or to believe are strongly influenced by the weight of reasons, and by how much of an effect those reasons have the potential to have on what we ought to do, in the context of the other reasons that there are. When there is a reason to believe something, but it is of low enough weight that it obviously makes no difference, given the other reasons in play in the situation, to what it is rational to believe, pragmatic considerations successfully predict that we will incorrectly find it intuitively compelling that there is in fact no reason to believe in such a situation.[5]

That pragmatic factors should affect our intuitive judgments about what there is no reason to do, is predicted by general Gricean considerations, in conjunction with the plain fact that reasons differ widely in their weight. Think of reasons like the considerations which go into the 'pros' and 'cons' columns that God might make in an exhaustive evaluation of what you should do. Since God's capacities are unlimited, He may be able to think of an awful lot of aspects of the situation which bear consideration in either the 'pros' column or the 'cons' column. We, however, are mere finite creatures – and ones who suffer from severe cognitive processing limitations, at that. We don't have time or energy for the kind of exhaustive search of which God is capable. We only have time to look at the top of God's lists. So long as we know that the reasons in favor that we have been able to survey are good *enough* to outweigh any comers, it is not that important to find out what those comers are – for we know in advance that they are insufficient anyway.

If any picture even roughly like this one is correct, and we are mostly interested in reasons in order to figure out what we ought to do or to believe, then reasons that are too far down God's list of pros and cons to be ones that it makes sense for us to bother with will be ones that it is not worth talking about. So this predicts, first, that pointing out or focusing on reasons of negligible or trivial weight will be conversationally unhelpful, at best. In fact, we should predict that even reasons of non-negligible weight will be not worth talking about, provided that they are obviously insufficient. And so this yields a second prediction: that pointing out or focusing on reasons that are obviously insufficient will be conversationally unhelpful, at best.

In fact, we can do more than make these two predictions. We can check to see that they are in fact borne out in cases. To see how, consider the case of Tom Grabit, familiar from the literature on

---

[5] Here I am summarizing arguments I have developed elsewhere in greater detail. For further discussion of the points which follow, see chapter five of Schroeder [2007] and further references there.

undercutting defeat.[6] In case 1, you see Tom come out of the library, pull a book out from under his shirt, cackle gleefully, and scurry off. You think, 'Hey – Tom just stole a book!' Intuitively, in this case you have some evidence for your conclusion – some evidence that Tom just stole a book. But now consider a twist on this case. Case 2 is the same as case 1, except that you know that Tom has an identical twin, Tim, from who you cannot visually distinguish him. Intuitively, in this case you don't have a reason for your conclusion after all – you don't have evidence that Tom stole a book. In case you lack this intuition, it is important to note that this is one of the paradigm cases of undercutting defeat in the literature, and undercutting defeater cases are often *defined* as cases in which, because of the defeater, you don't have any reason for your conclusion after all.[7] So even if you don't share this intuitive judgment about case 2, it is an intuitive judgment that has been widely shared by epistemologists.

But now compare case 3, which is like case 2, except that now you know that Tom and Tim have a third identical sibling, Tam, from whom you can visually distinguish neither. I submit that in case 3, you have even less reason to believe that Tom stole a book than in case 2, and hence that you must have had some reason to believe that Tom stole a book even in case 2. If this is unobvious, observe that it would take more additional evidence that Tom stole a book to make you justified in drawing the conclusion that Tom stole a book in case 3 than in case 2. If this is right, and you really do have a reason to believe that Tom stole a book in case 2, then the intuitive judgment about that case that has been shared by so many epistemologists over the years is mistaken. And that is precisely what was predicted by our pragmatic considerations: we predicted that reasons of relatively low weight would seem not to be reasons at all.

In fact, by constructing a series of Grabit cases, we can show that, just as our pragmatic considerations predicted, our intuitive judgments can be misleading even for reasons of non-trivial weight, provided that they are obviously outweighed. Case 4, in which you know Tom to be only one of four identical siblings, shows that your reason to believe that Tom stole a book didn't go away in case 3, either, but was merely reduced in weight. And case 5, in which you know that he is only one of five, shows that your reason didn't go away in case 4, either. By such reasoning, we can construct a sequence of reasons of descending weight. In such a sequence, the reason in case 2 is relatively weighty! But as decades of epistemology have shown, philosophers still systematically have shared the intuition that there was no reason even in case 2. This is easy to explain, and in fact exactly what our pragmatic considerations predict, because your reason is obviously insufficient. In the case described, it is obvious that the thing for

---

[6] See Lehrer and Paxson [1969].
[7] Compare, for example, Pollock and Cruz [1999, 195-197].

you to do is to withhold judgment about whether Tom stole a book, whether you have a reason to believe it or not.

So the Grabit cases, because they are subject to an independent argument that our intuitive judgments are wrong, provide evidence that the pragmatic effects that we earlier predicted based on abstract considerations are genuinely in play, and really do affect the way that we think about philosophical examples. The idea that our intuitive judgments about what someone does or does not have reason to do might be affected by pragmatic factors is therefore not just an abstract possibility. It is a very real phenomenon that we need to be careful about in evaluating intuitive judgments about reasons or evidence. Consequently, it safer to go by our indirect evidence. If unjustified beliefs do not improve our evidential position for their consequences, then even if we do have reasons to believe their consequences, those reasons must be of little or no weight, and consequently will be ones about which we will have misleading negative intuitions. And this is precisely what the argument of the last section predicts. It predicts that any reasons we have by having unjustified beliefs will be utterly defeated – totally reduced in weight.

5 the payoff for epistemology?

So far in this paper, I have been concentrating my energies on rebutting two arguments for an apparently abstruse conclusion: that there must be a relatively high bar on the 'having' relation that is involved in having evidence. Even if my rebuttal is successful, one might wonder about its point. The point of all of this hard work, I claim, is that it expands the options in epistemology in a fruitful way, which makes it easier to understand how basic perceptual beliefs could be based on evidence. This is because the assumption of a high bar on what it takes to have evidence – despite its overwhelming initial plausibility – has been complicating the dialectic about basic perceptual epistemology, and consequently the dialectic about foundationalism and coherentism, internalism and externalism, and rationalism and empiricism, for a very long time.

To see why, think about our high bar as at bottom a high bar on what it takes to have evidence in cases of *inferential* justification – cases in which one belief is based on evidence you have in virtue of having another belief. Inferentially justified beliefs can meet this bar, provided that they are based on other beliefs that are justified. So inferentially justified beliefs can be based on evidence. But the problem is whether *non*-inferentially justified, *basic* perceptual beliefs can also be based on evidence, on this view. And the problem essentially comes in two parts. The first part of the problem is that the high bar on having evidence for inferentially justified beliefs looks like it is too high to be met, in the case of basic perceptual

13

beliefs, if those are justified only on the basis of perceptual experiences. So that means that if basic perceptual beliefs are to count as based on evidence at all, then in order to keep the high bar on what it takes to have evidence for inferentially justified beliefs, we need to adopt a *mixed* view, with mixed standards about what it takes to have evidence – sometimes requiring a justification condition, and sometimes not. But the second part of the problem is that this kind of view looks *ad hoc*, unstable, and explanatorily unsatisfactory.

In the remainder of this section I'll rehearse this kind of dialectic in more detail; my goal will not be to definitively set aside any of the options that come up along the way, but only to emphasize that dropping the high bar on what it takes to have evidence relaxes the constraints on perceptual epistemology in an important way. It makes it much easier to see how to tell a unified explanatory story about how beliefs can be based on evidence, in perceptual and non-perceptual cases alike. Perhaps this view will suffer from its own problems, as well, but that should be a matter for theorizing to decide, not to be stipulated in advance by ruling it out, out of court.

To take the first side of the problem first, let's look at whether basic perceptual beliefs could be based on evidence, if the high bar has to be satisfied. So to address that, take first the thesis that H=JB, which was one kind of view which upheld the high bar. We are assuming that a belief is based on evidence only if it is based on evidence that you have. So if having evidence requires justifiedly believing it, it follows from this view that beliefs can be based on evidence only if they are based on other beliefs. That is, non-inferentially justified beliefs can't be based on evidence at all. But this is a highly counterintuitive result! It is one thing to say that some beliefs are justified without being justified on the basis of other beliefs. Any foundationalist accepts this. It is another thing to say that basic perceptual beliefs are not based on any evidence at all. This is a bizarre claim – beliefs formed directly on the basis of perceptual experiences are intuitively a *paradigm* of beliefs that are based on evidence, because perceptual experience has long been a paradigm method of gaining evidence about the world. Now, it is easy to see why a theory might lead some philosophers to endorse the conclusion that this is misleading, and basic perceptual beliefs are not, after all, based on any evidence at all. But it smells to me like at least the sort that we should shy away from without some clear motivation. So it's worth looking at how we can avoid having to draw this conclusion.

So far, the problem arises from the *B* part of H=JB, rather than the *J* part. But the high bar focuses on the *J* part – the idea that there must be – at least – a justification condition on having evidence. So one thought *might* be to generalize. Maybe perceptual experiences, like beliefs, can be ways of having

14

evidence – provided that *they* are justified. This leads to the identification, H=JPres[8], which preserves the justification as a minimal bar for having evidence, but is liberal about the kind of *attitude* that having a reason can consist in – it can consist in belief in cases of inferential justification, and perceptual experiences in cases of basic perceptual justification.

The only problem with this view is that perceptual experiences aren't the kind of thing that *can* be justified. Beliefs can be justified or unjustified because they fall under the reach of rational criticism, and they fall under the reach of rational criticism because they are states that we hold for reasons. Perceptual states are not things we enter into for reasons. True, we can choose whether to look in the direction of the stick in the water or to look away. But once we look in its direction, its looking bent to us is not something that we can control, it is not something that we do for reasons, it is not subject to rational criticism, and it is not subject to justification. Don't take my word for it; Sellarsians tell us that it is 'outside the space of reasons'. All we need for our problem here is that perceptual experiential states (as opposed to perceptual judgments, which are based on perceptual states) do not themselves admit of justification. If not, then H=JPres is a failed venture.

One way of trying to get the benefits of H=JPres without its costs is to double down on the high bar by endorsing H=K. Maybe perceptual states do not themselves admit of justification, this view goes, but they *do* admit of constituting *knowledge*. Take any of our wide range of factive experiential verbs, like 'sees that', 'observed that', and 'smelled that'. Williamson [2000] tells us that each of these implies 'knows that'. If some perceptual states actually constitute knowledge, then by H=K, these perceptual states will constitute ways of having evidence on which beliefs can be based. Hence, non-inferentially justified beliefs can be based on evidence – not evidence provided by other beliefs, but evidence provided by perceptual states. On this view, all justified beliefs are based on evidence; some are based on evidence provided by other beliefs which constitute knowledge, and that is the inferential case, whereas others are based on evidence provided by perceptual states which constitute knowledge, and that is the basic case. But in both cases, on this view, all justified beliefs are based on evidence.

That all sounds very good (though there are important and not-to-be-glossed-over complications in the details[9]). But it has one very striking problematic feature. And that is that on this picture, there is a

---

[8] Here as before, think of 'Pres' as just a placeholder for some genus of attitude which is general enough to include both belief and perceptual experiences as special cases. I think of both as 'presentational attitudes', in the sense that each presents its content to its holder *as true* of the actual world. In this way, I think, belief and perception differ from other cognitive attitudes like wondering, assuming, supposing, and imagining.

[9] For example, according to Williamson [2000], knowledge entails belief. So on that view, the perceptual states that constitute knowledge would also constitute belief. Since they would not be based on further evidence, they would therefore be knowledge

very important difference between veridical and non-veridical perceptual experiences' abilities to play a role in justification. This is because even if some perceptual experiences constitute knowledge, non-veridical perceptual experiences certainly don't – for knowledge requires truth. That means that this view predicts that internal duplicates with duplicate pasts in very similar environments can nevertheless differ not just in whether they know, but in whether their belief is justified, if for the very first time one sees something green in front of her while the other experiences a vivid hallucination as of something green in front of her, and both form the perceptual belief that there is something green in front of them. This is extremely unintuitive.[10] Like the thesis that some beliefs are justified without being based on any evidence itself, it is the kind of thesis one can imagine talking oneself into, if one has the right kind of theory. But again, it smells like the kind of conclusion that we should shy away from, if we can.

So unfortunately the high bar on having evidence is not a bar that can be met in the full range of basic perceptual cases. At most it can be met in the case of veridical perceptual experiences, but it would be best to avoid needing to distinguish between the justification provided by veridical and by non-veridical perceptual experiences. So if we want to also both avoid the conclusion that some beliefs are justified on the basis of no evidence whatsoever, and hang onto the high bar in the cases of inferentially justified belief appealed to by the truism, there is just one way to go: *mixed views*: views which propose high standards on having evidence in the cases of inferential justification, and low standards in the case of basic perceptual justification.

This is what a variety of epistemologists try to do. On views like these, there are *two ways* to 'have' evidence. In order to have some kinds of evidence, you have to justifiedly believe it. For example, to have Max's smiling as evidence that Max is happy, you have to justifiedly believe that Max is smiling. But according to this view, for other kinds of evidence, only something much weaker is required. In fact, according to many of these views, you can have some kinds of evidence just by virtue of its being true.[11]

---

that was not based on any evidence. Williamson avoids the question of whether all knowledge is based on evidence, by observing that on his view, all knowledge is *supported* by the evidence – because if you know that *p*, then *p* is among your evidence, and that supports *p*. This simply evades the question of whether someone's knowledge that *p* can be *based* on evidence.

[10] Note that what I am claiming is unintuitive here is not just a feature of *internal duplicates*. Some philosophers believe that internal duplicates must either both be justified or neither be justified in their same beliefs, but I am not relying on that here. For example, internal duplicates might have different pasts, or be differentially reliable, or be in different environments, one more conducive to truth belief formation than the other, and many philosophers believe that these conditions can affect whether they are justified or not. What I am pointing out here, is that two people might be internal duplicates, with identical pasts, equally reliable, in environments which are identical up to the very thing that each is looking at right now, which one veridically sees to be green and the other experiences a perceptual illusion as of being green. Surely in *this* case one is justified in her perceptual belief formed on the basis of this perceptual experience if the other is. But the H=K thesis only admits, at best, veridical perceptual experiences as among their evidence.

[11] Compare, for example, Pryor [2005] and Conee and Feldman [2004].

For example, to have your perceptual experience as of something green in front of you as evidence that there is something green in front of you, you need only to have a perceptual experience as of something green in front of you – you don't need to have any belief about this perceptual experience at all – let alone a justified one.

Now, I ultimately want to agree that *when* you have such a perceptual experience you have evidence for the belief that there is something green in front of you – even if you have no prior beliefs about your perceptual experience, let alone justified ones. So I think there is something descriptively accurate about this view – in contrast to views which hold that since 'having' evidence requires having a prior belief, your belief that there is something green in front of you can be justified only if you first believe *that you perceive* as of something green being in front of you.[12] But there are nevertheless at least three serious problems with mixed views.

The first problem is that it is easy to provide an error theory for what could lead proponents of mixed views to hold that you have certain sorts of evidence simply in virtue of its being true. Certainly, given that your perceptual experiences are even moderately reliable, the fact that you are having a perceptual experience as of something green in front of you is evidence in the *objective* sense that there is something green in front of you. It is objective evidence even for people who are not you and do not know about your perceptual experience. This makes the move of calling this perceptual experience part of your evidence even though you don't satisfy the epistemic access ('having') condition for subjective evidence very suspicious – after all, it is *true* that it is part of your evidence in *one* sense of 'evidence'. It is part of the objective evidence. What is at stake is whether it is part of your *subjective* evidence. But in all of the non-basic cases, even this view admits that some further condition is required.

The second problem for mixed views like this one is larger. If the only evidence you have that there is something green in front of you is that it perceptually appears to you as if there is something green in front of you, then there is a serious question of why this evidence is so good. It is a big jump, after all, to get from 'it perceptually appears to me as if *p*' to '*p*'. In fact, it is precisely this jump that plays a central role in some important kinds of skeptical argument, which point out that it is consistent with all of your

---

[12] The idea that having evidence involves a *having* relation makes it easy to see why early 20[th]-century empiricists were so wedded to the idea that the epistemologically basic beliefs were existential beliefs about perceptual experiences, and hence that you need to have beliefs *about* your perceptual experiences, in order for them to play a role in justifying your other beliefs. For example, compare Lewis [1946] and Chisholm [1957]. But not only was this view descriptively inaccurate, incorrectly classifying subjects who lacked beliefs about their own perceptual experiences as being unable to have beliefs justified on the basis of perception, they ran into troubles both explaining how beliefs about perceptual experiences could be justified in the first place, and how beliefs about perceptual experiences could provide as good of evidence about the external world as perceptual experience seems to provide. Compare Pollock [1974].

evidence that you are a brain in a vat. Whatever the correct response to other kinds of skeptical arguments, whether this particular skeptical argument gets going depends on whether your evidence includes only facts about your own perceptual state, or includes information about the world. Even putting such skeptical arguments aside, it is hard to see what makes perception an especially good source of evidence about the external world, if the only evidence it gives us directly is about our own mind. Even if you can have the fact that it perceptually seems to you that *p* as evidence just by virtue of its being true – and thus without needing to stand in any further positive epistemic relation to it, such as believing or perceiving it to be the case – it is at least not obvious why this is particularly good evidence to have.

I point this out not in order to argue that this is an unsolvable problem for views according to which your basic perceptual evidence about the external world consists in facts about your own perceptual psychology. Certainly many philosophers have tried to wade into this breach and explain why this is nevertheless good evidence about the external world. But my point is that it is a problem that has a particular kind of *source*, which is illustrated by how easily the contrasting views avoid this pitfall. According to the contrasting picture, in its perceptually seeming to you as if *p*, you have *p* as evidence that *p* (and count as 'having' it because you stand in a positive epistemic relation to it, to boot). On this sort of picture, it is easy to see why that is a good situation to be in, and hence why perception is a privileged source of evidence about the external world. This is because in its perceptually seeming to you that *p*, you come to have *p* as a reason to believe that *p* – and no other evidence that *p* could be better.

Finally, the third problem for this kind of view is I think quite serious. If there are two such different ways of 'having' evidence – such that one way requires that you justifiably believe the evidence, and the other requires that the evidence be true – then there will be two very different ways of *basing* a belief on evidence. But that calls into question what, exactly, this sort of theorist has gained by saying that all justified beliefs are based on evidence. They are based on evidence in such different ways, that it stretches the imagination to understand how there is even a unified sense of 'based on the evidence' in which both are based on the evidence.

Of course, as theorists we can *use* the words, 'based on the evidence' however we like. We can even stipulate that no matter how beliefs are justified, we will say that they are 'based on the evidence'. But that falls far short of it turning out that there is such a thing as being based on the evidence which happens with inferentially justified beliefs, and which also happens – the very same thing – with basic perceptual beliefs. Again, the idea that there is no single relation of being based on the evidence, such that holding that sense of 'based on evidence' fixed, all justified beliefs are based on evidence, is the kind of thing that a theory can

lead one to accept. But my sense of smell is getting to me again, and it tells me that this is the kind of thesis that we should avoid, if there is any easy way to do so.

It is crucially important not to underestimate the centrality of this problem in epistemology. Repeatedly in the history of twentieth-century epistemology, philosophers have returned to the idea that there are simply different things going on in the case of basic perceptual justification and in the case of inferential justification – that one philosophical story is required for one, and another quite different story for the other. Mixed views fall squarely within this tradition. And the prevalence of mixed views like this has been one of the strongest original motivations for both coherentism and externalism. Coherentists can offer a unified account of justification, by applying what goes for the inferential case to what are apparently basic perceptual cases.[13] And since the beginning of externalism in epistemology, externalists have argued that foundationalists have to accept externalist explanations of basic perceptual justification anyway, and that once you take that on board, they are merely extending what foundationalists accept anyway to the inferential cases.[14] It's hard to complain about causal or pure reliabilist theories of knowledge, when your own story about basic perceptual justification is distinguishable from them only by being less unified and less explicit.[15] So for these reasons I take it to be quite a serious charge that the view we've been considering does not provide a unified picture of having evidence and basing beliefs on evidence. It would be nice to be able to do better.

6   wrap-up

Just to rehearse: in the last two sections we have encountered a familiar dialectic. The arguments I've sought to rebut in this paper would have us place a high bar on what it takes to have evidence for inferentially justified beliefs. That high bar cannot be met by basic perceptual beliefs, which are based on perceptual experiences, rather than on beliefs. Since perceptual experiences don't admit of justification, it doesn't help to liberalize H=JB to H=JPres. And since knowledge requires truth, the strategy of doubling down on H=K only nets cases of veridical perceptual experiences, even if its details can be properly worked out.

So if the high bar applies across the board, then we are led to the conclusion that basic perceptual beliefs cannot be based on any evidence. In order to retain the high bar on having evidence for inferentially

---

[13] See especially Bonjour [1985].
[14] See especially Armstrong [1973].
[15] On this instability of mixed views, compare McDowell [1994].

justified beliefs, but make it possible for basic perceptual beliefs to be based on evidence, some have advocated mixed views. But mixed views appear to conflate objective and subjective evidence, make it puzzling why the evidence provided by perceptual experiences is particularly good and open at least one door to the skeptic, and present a problematically disunified picture of justification that is vulnerable to attack from both sides.

In the context of this dialectic, I *do* think it is worth taking a closer look at the motivation for raising the bar on having evidence so high in the inferential case. By lowering that bar, we can endorse H=Pres, and thereby offer a single, unified account of what it is to have evidence that is satisfiable in both inferential and basic cases, without excluding non-veridical perceptual states. This unified account is very simple: it says that a belief is propositionally justified just in case it is supported by the balance of the undefeated evidence the agent has, and it understands having evidence as everywhere the same: a matter of bearing some presentational attitude (either belief or a perceptual state) toward a proposition.

Moreover, because it is a unified account[16], we can better fend off the challenges from coherentists, on the right, and from pure externalists, on the left, who seek to generalize either our inferential case or our basic case to include all of justification. I also believe that such a unified account promising to be more deeply explanatory, and even to tell us not just *when* beliefs are justified, but what justification *is*. I happen to find this picture attractive because it avoids all of our earlier pitfalls. But all I have been trying to argue for here, is that it is a serious alternative, rather than one that is ruled out of court by a truism that we all need must recognize. If something is wrong with it, that is for the theoretical balancing of costs and benefits to decide, not a snap judgment that it doesn't merit consideration, due to the truism.[17]

---

[16] Whether this is, in fact, a unified account depends on whether there is anything unified about the category, *Pres*, which includes both beliefs and perceptual states. I am optimistic on this score, but making good on this is one of several undischarged obligations of the view I've barely sketched here.

[17] Special thanks to Jake Ross, as well as to conversations with Jim Pryor and Timothy Williamson. And thanks to Andrew Reisner and Asbjørn Steglich-Peterson.

references

Armstrong, David [1973]. *Belief, Truth, and Knowledge*. Cambridge: Cambridge University Press.

Bonjour, Laurence [1985]. *The Structure of Empirical Knowledge*. Cambridge, MA: Harvard University Press.

Chisholm, Roderick [1957]. *Perceiving: A Philosophical Study*. Ithaca, NY: Cornell University Press.

Conee, Earl, and Richard Feldman [2004]. *Evidentialism*. Oxford: Oxford University Press.

Feldman, Richard [1988]. 'Having Evidence.' Reprinted in Conee and Feldman [2004], 219-241.

Lehrer, Keith, and Thomas Paxson [1969]. 'Knowledge: Undefeated Justified True Belief.' *The Journal of Philosophy* 66(4): 225-237.

Lewis, C.I. [1946]. *An Analysis of Knowledge and Valuation*. LaSalle, IL: Open Court.

McDowell, John [1994]. *Mind and World*. Oxford: Oxford University Press.

Pollock, John [1974]. *Knowledge and Justification*. Princeton: Princeton University Press.

Pollock, John, and Joseph Cruz [1999]. *Contemporary Theories of Knowledge*, second edition. New York: Rowman and Littlefield.

Pryor, James [2005]. 'There is Immediate Justification.' In Matthias Steup and Ernest Sosa, eds., *Contemporary Debates in Epistemology*. Oxford: Blackwell.

Schroeder, Mark [2007]. *Slaves of the Passions*. Oxford: Oxford University Press.

_____ [2008]. 'Having Reasons.' *Philosophical Studies* 139(1): 57-71.

_____ [unpublished]. 'Stakes, Withholding, and the Subject-Sensitivity of 'Knows'.'

Williamson, Timothy [2000]. *Knowledge and Its Limits*. Oxford: Oxford University Press.